



并行与分布式计算

Parallel & Distributed Computing

陈鹏飞
数据科学与计算机学院
2018-03-23



Lecture 3 — Parallel Programming Model

Pengfei Chen

School of Data and Computer Science

March 23, 2018

Outline:

1

Introduction

2

Share Memory Model

3

Message Passing Model

4

GPGPU Programming Model

5

Data Intensive Computing Model




Parallel programming models

INTRODUCTION

Examples

➤ Computing π

$$\left. \begin{array}{l} \arctan(1) = \pi/4 \\ \arctan(0) = 0 \\ \frac{d}{dx} \arctan(x) = 1/(1+x^2) \end{array} \right\} \Rightarrow \pi = \int_0^1 \frac{4}{1+x^2} dx$$






Examples

➤ Computing π : Sequential Code

```
double compute_pi(int n) {  
    double sum = 0.0;  
    for (int i = 0; i < n; i++) {  
        double x = (i + 0.5) / n;  
        sum += 1.0 / (1.0 + x*x);  
    }  
    double pi = 4.0 * sum / n;  
    return pi;  
}
```



Examples


➤ Computing π : *POSIX Thread Version*

```
// global variables
int thread_count;
int n;
double* local_sum;

// multithreaded version
double compute_pi () {
    thread_handles = (pthread_t*) malloc (thread_count*sizeof(pthread_t));
    local_sum = (double*) malloc (thread_count*sizeof(double));
    // parallel part
    for (thread = 0; thread < thread_count; thread++)
        pthread_create(&thread_handles[thread], NULL, thread_sum, (void*)thread);
    for (thread = 0; thread < thread_count; thread++)
        pthread_join(thread_handles[thread], NULL);
    // sequential part
    double sum = 0.0;
    for (thread = 0; thread < thread_count; thread++)
        sum += local_sum[thread];
    double pi = 4.0 * sum / n;
    return pi;
}
```



```
void* thread_sum(void* rank) {
    int my_rank = (int) rank;
    double my_sum = 0.0;
    // domain decomposition
    for (int i = my_rank; i < n; i += thread_count) {
        double x = (i + 0.5) / n;
        my_sum += 1.0 / (1.0 + x * x);
    }
    local_sum[my_rank] = my_sum;
    return NULL;
}
```




A. Grama et al., "Introduction to Parallel Computing," Addison Wesley, 2003

Examples

➤ Computing π : *OpenMP Version*

```
/* compile as $> gcc -fopenmp -lm  
* run as      $> OMP_NUM_THREADS=2 ./a.out  
*/
```

```
double compute_pi(int n) {  
    int i;  
    double sum = 0.0;  
    #pragma omp parallel for reduction(+: sum) schedule(static)  
    for (i = 0; i < n; ++i) {  
        double x = (i + 0.5) / n;  
        sum += 1.0 / (1.0 + x * x);  
    }  
    double pi = 4.0 * sum / n;  
    return pi;  
}
```



Examples

➤ Computing π : *MPI Version*

```
#include "mpi.h"
#include <stdio.h>
#include <math.h>
int main( int argc, char *argv[] ) {
    int n, my_rank, numprocs, i;
    double mypi, pi, h, sum, x;
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD,&my_rank);
    while (1) {
        if (my_rank == 0) {
            printf("Enter the number of intervals: (0 quits) ");
            scanf("%d",&n);
        }
        MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
        if (n == 0) break;
        else {
            sum = 0.0;
            for (i = my_rank; i < i += numprocs) {
                x = (i + 0.5) / n;
                sum += 4.0 / (1.0 + x*x);
            }
            mypi = sum / n;
            MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0,
                MPI_COMM_WORLD);
            if (my_rank == 0) printf("pi is approximately %.16f\n", pi));
        }
    }
    MPI_Finalize();
    return 0;
}
```



Examples

➤ Computing π : OpenCL Version

/ to be executed on device */*

```
__kernel void pi(const int niters, const float step_size,
__local float* local_sums, __global float* partial_sums) {
    int num_wrk_items = get_local_size(0);
    int local_id = get_local_id(0);
    int group_id = get_group_id(0);
    float x, accum = 0.0f;
    int i, istart, iend;
    istart = (group_id * num_wrk_items + local_id) * niters;
    iend = istart + niters;
    for(i = istart; i < iend; i++) {
        x = (i + 0.5f) * step_size;
        accum += 4.0f / (1.0f + x * x);
    }
    local_sums[local_id] = accum;
    barrier(CLK_LOCAL_MEM_FENCE);
    reduce(local_sums, partial_sums);
}
```

/ to be executed on host */*

```
pi_res = 0.0f;
for (unsigned int i = 0; i < nwork_groups; i++) { pi_res += h_psum[i]; }
pi_res *= step_size;
```

Examples

➤ Computing π : *Summary*

Different parallel granularity

Different parallel implementation (implicit VS explicit)

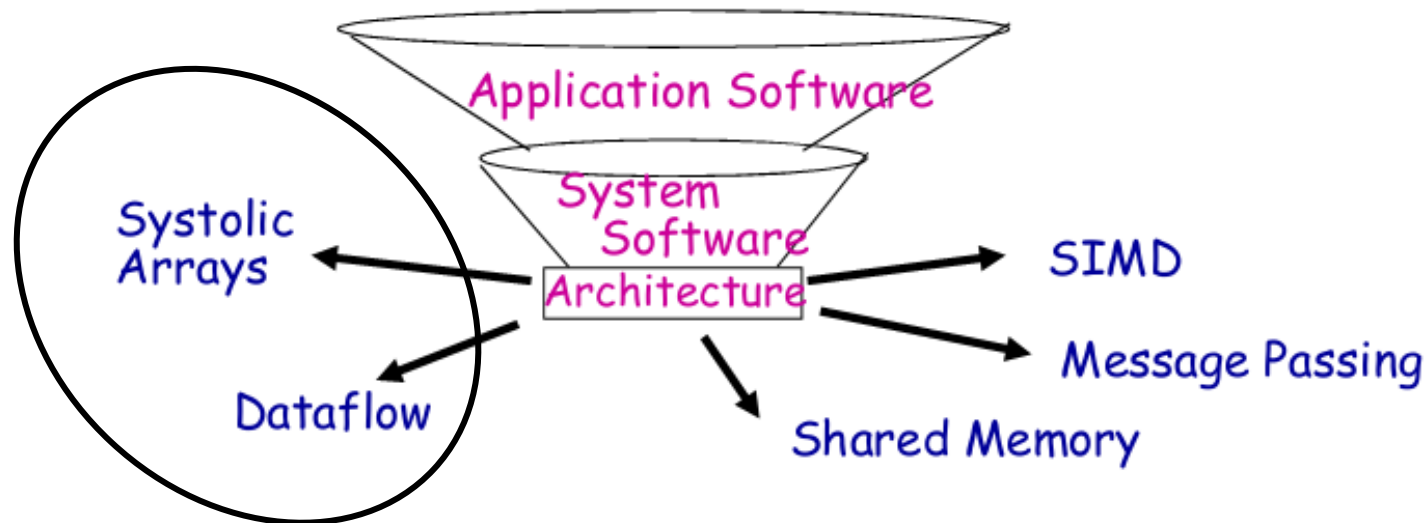
Different code scale

Different hardware architecture

Different what?

History

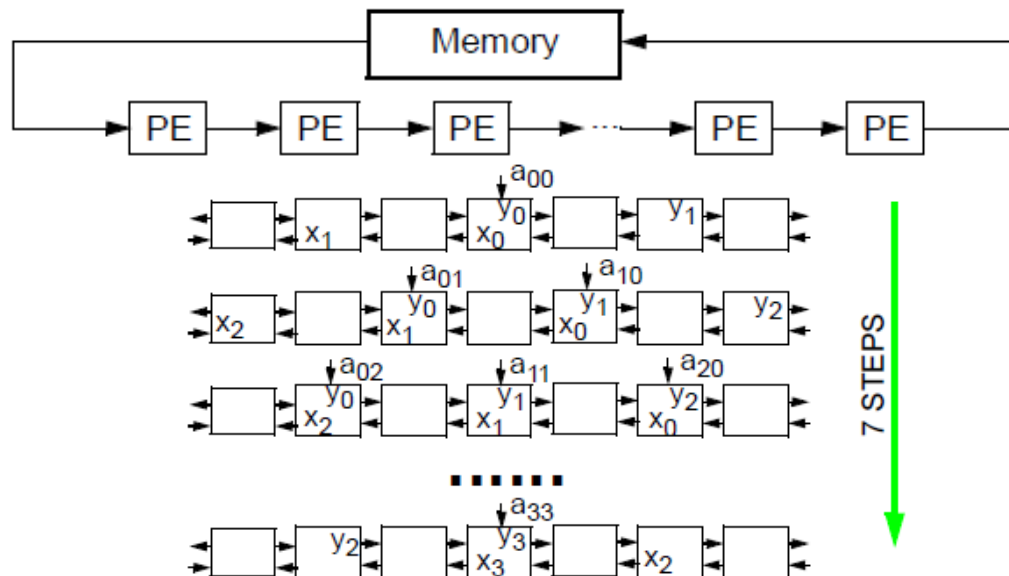
- Historically (1970s – early 1990s), each parallel machine was unique, along with its programming model and language
 - Architecture = prog. model + comm. abstraction + machine;
 - parallel architectures tied to programming models;
- Divergent architectures, with no predictable pattern of growth



Uncertainty of direction paralyzed parallel software development!

Systolic Array

- Types of special-purpose computers:
- Inflexible and highly dedicated structures
 - Structures, enabling some programmability and reconfiguration.
 - Data would move through the system at regular “heartbeats” as determined by local state.



Inner product step (ISP) cell:

$$y_{out} = y_{in} + x_{in} \times a_{in}$$

$$x_{out} = x_{in}$$

Systolic vector-matrix multiplication



Today

- Nowadays we separate the **programming model** from the underlying **parallel machine architecture**
 - Dominant: shared address space, message passing, data parallel
 - Others: data flow, systolic arrays
- Extension of “computer architecture” to support communication and cooperation
 - OLD: Instruction Set Architecture
 - NEW: Communication Architecture
- Defines
 - Critical abstractions, boundaries, and primitives (interfaces)
 - Organizational structures that implement interfaces (hw or sw)
- **Compilers, libraries and OS are important bridges**



Programming Model

- What programmer uses in coding applications
- **Specifies communication and synchronization**
- **Examples**
 - Multiprogramming: no communication or synch. at program level
 - Shared address space: like bulletin board
 - Message passing: like letters or phone calls, explicit point to point
 - Data parallel: more strict, global actions on data
 - Implemented with shared address space or message passing



Programming Model

➤ von Neumann model

- Execute a stream of instructions (machine code)
- Instructions can specify
 - Arithmetic operations
 - Data addresses
 - Next instruction to execute
- Complexity
 - Track billions of data locations and millions of instructions
 - Manage with
 - ✓ Modular design
 - ✓ High-level programming languages (isomorphic)

Programming Model

➤ Parallel Programming Models

□ Message passing

- Independent tasks encapsulating local data
- Tasks interact by exchanging messages

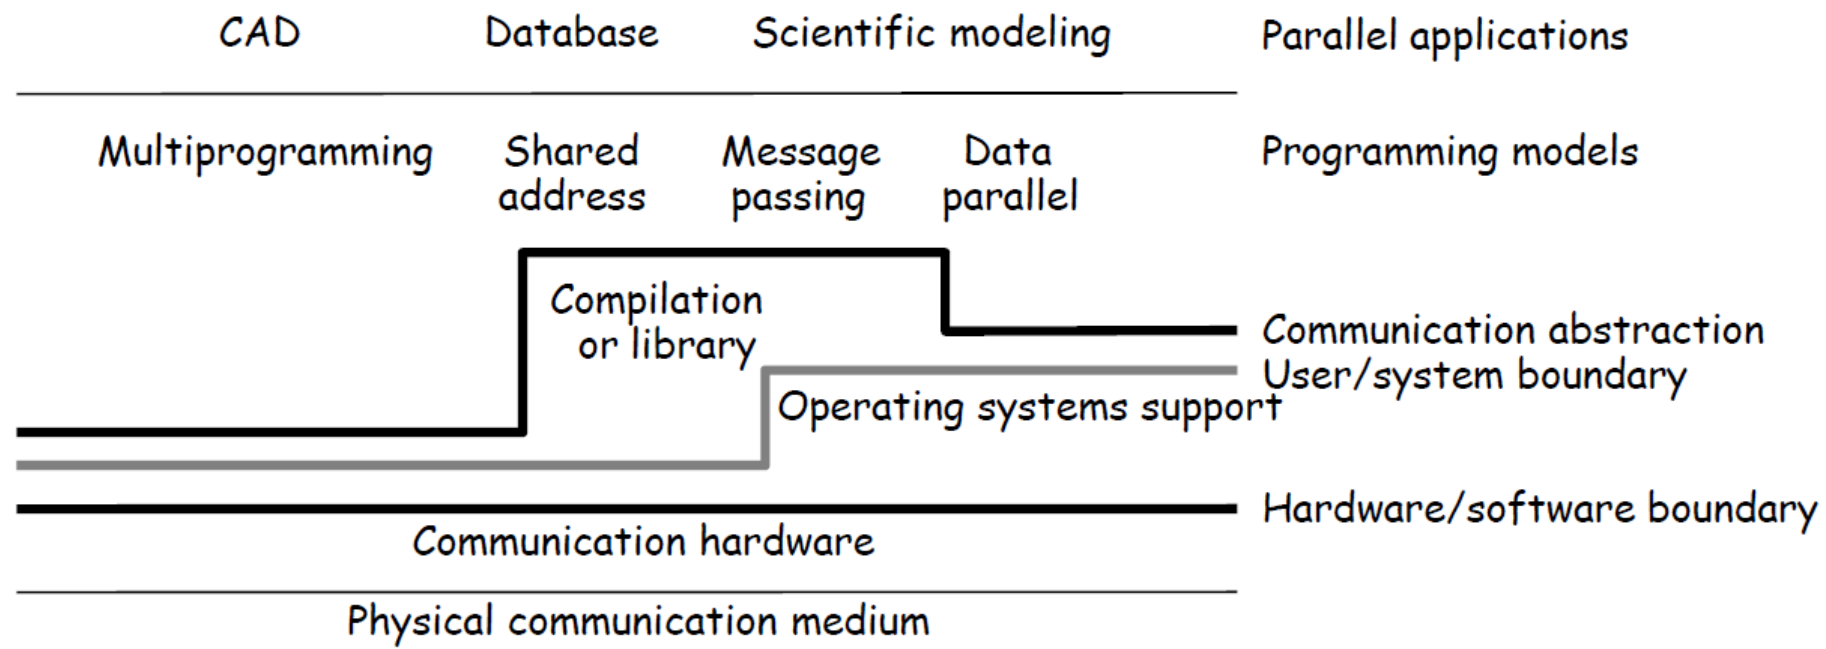
□ Shared memory

- Tasks share a common address space
- Tasks interact by reading and writing this space asynchronously

□ Data parallelization

- Tasks execute a sequence of independent operations
- Data usually evenly partitioned across tasks
- Also referred to as “embarrassingly parallel”

Modern Layered Framework

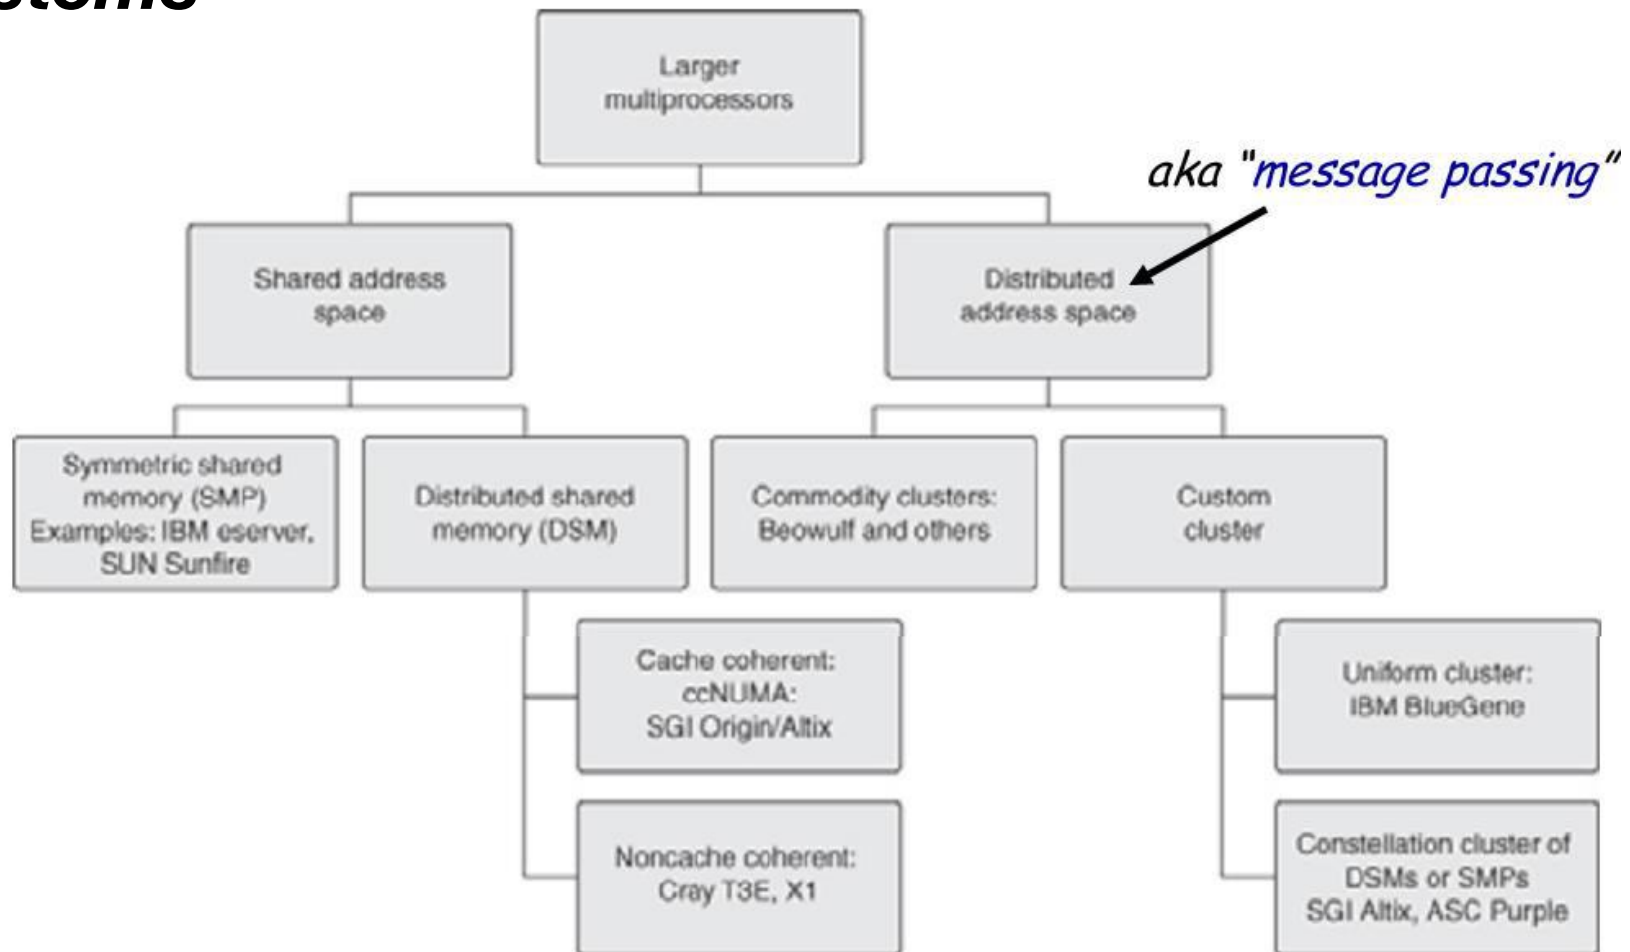




Evolution of Architectural Models

- **Historically, machines tailored to programming models**
 - ❑ Programming model, communication abstraction, and machine organization lumped together as the “architecture”
- **Evolution helps understand convergence**
 - ❑ Identify core concepts
- **Most common models**
 - ❑ Shared memory model, threads model, distributed memory model, GPGPU programming model, data intensive computing model
- **Other models**
 - ❑ Dataflow, Systolic arrays
- Examine programming model, motivation, intended applications, and contributions to convergence

Taxonomy of Common Large-Scale SAS and MP Systems



© 2007 Elsevier Inc. All rights reserved.

Aspects of a Parallel Programming Model

➤ Control

- ❑ How is parallelism created?
- ❑ In what order should operations take place?
- ❑ How are different threads of control synchronized?

➤ Naming

- ❑ What data is private vs. shared?
- ❑ How is shared data accessed?

➤ Operations

- ❑ What operations are atomic?

➤ Cost

- ❑ How do we account for the cost of operations?



Parallel programming models

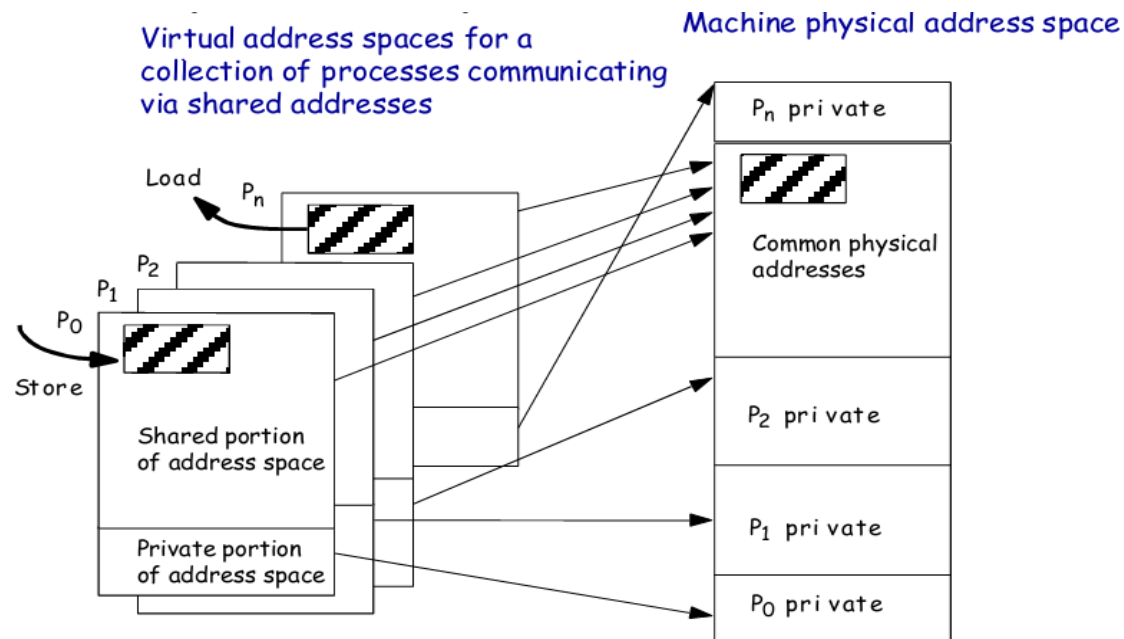
SHARED MEMORY MODEL

Shared Memory Model

- Any processor can **directly** reference any memory location
 - ❑ Communication occurs implicitly as result of loads and stores
- Convenient
 - ❑ Location transparency
 - ❑ Similar programming model to time-sharing on uniprocessors
 - Except processes run on different processors
 - Good throughput on multiprogrammed workloads
- Popularly known as *shared memory* machines or model
 - ❑ Ambiguous: memory may be physically distributed among processors

Shared Memory Model

- Process: virtual address space plus one or more threads of control
- Portions of address spaces of processes are shared



- Writes to shared address visible to other threads, processes
- Natural extension of uniprocessor model: conventional memory operations for comm.; special atomic operations for synchronization

Shared Memory Model

- In this programming model, **tasks share a common address space**, which they read and write asynchronously 异步
- Various mechanisms such as locks / semaphores may be used to control access to the shared memory 使用锁或者信号来完成对共享内存的访问
- An advantage of this model from the programmer's point of view is that the notion of data "ownership" is lacking, so there is **no need to specify explicitly the communication of data between tasks**
 - ❑ Program development can often be simplified 明确

Shared Memory Model

- An important disadvantage in terms of performance is that it becomes more difficult to understand and manage **data**

locality

数据局部性

- Keeping data local to the processor that works on it **conserves**

节省内存访问

memory accesses, cache refreshes and bus traffic that occurs

when multiple processors use the same data

在处理器本地储存数据

- Unfortunately, controlling data locality is hard to understand and beyond the control of the average user



Shared Memory Model

➤ An important disadvantage in terms of performance is that it becomes more difficult to understand and manage **data**

locality

- ❑ Keeping data local to the processor that works on it conserves memory accesses, cache refreshes and bus traffic that occurs when multiple processors use the same data

- ❑ Unfortunately, controlling data locality is hard to understand and beyond the control of the average user

Implementations

- Native compilers and/or hardware translate user program variables into actual memory addresses, which are global
 - ❑ On stand-alone SMP machines, this is straightforward
- On distributed shared memory machines, such as the SGI Origin, memory is physically distributed across a network of machines, but made global through specialized hardware and software

SAS Machine Architecture

➤ One representative architecture: SMP

□ Used to mean *Symmetric MultiProcessor*


- All CPUs had equal capabilities in every area, e.g., in terms of I/O as well as memory access

□ Evolved to mean *Shared Memory Processor*

- Non-message-passing machines (included crossbar as well as bus based systems)

□ Now it tends to refer to *bus-based shared memory machines*

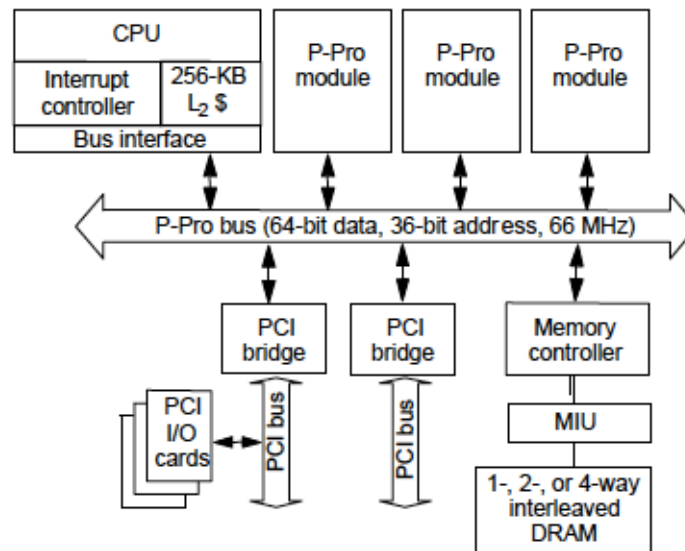

- Small scale: < 32 processors typically



对称多处理器




SAS Example: Intel Pentium Pro Quad




- All coherence and multiprocessing glue in processor module
- Highly integrated, targeted at high volume
- Low latency and high bandwidth

SAS Example: Sun SPARC Enterprise M9000




- 64 SPARC64 VII+ quad-core processors (i.e. 256 cores)
- Crossbar bandwidth: 245 GB/sec (snoop bandwidth)
- Memory latency: 437-532 nsec (i.e. 1050-1277 cycles @ 2.4 GHz)
- Higher bandwidth, but also higher latency

Scaling Up



“Dance hall”



Distributed memory


- Problem is interconnect: cost (crossbar) or bandwidth (bus)
- Dance-hall: bandwidth is not scalable, but lower cost than crossbar
 - ❑ Latencies to memory uniform, but **uniformly large**
- Distributed memory or non-uniform memory access (**NUMA**)
 - ❑ Construct shared address space out of simple message transactions across a general-purpose network (e.g. read-request, read-response)



Example: SGI Altix UV 1000




Blacklight at the PSC (4096 cores)



256 socket (2048 core) fat-tree
(this size is doubled in Blacklight via a torus)

- Scales up to **131,072 cores**
- **15GB/sec** links
- Hardware cache coherence



8x8 torus



Parallel programming models

THREAD MODEL



Threads Model


- This programming model **is a type of shared memory programming**
- In the threads model of parallel programming, a single process can have multiple, concurrent execution paths
- Perhaps *the most simple analogy* that can be used to describe threads is the concept of a single program that includes a number of subroutines

Threads Model

- Program is a collection of threads of control
 - ❑ Can be created dynamically, mid-execution, in some languages
- Each thread has a set of private variables, e.g., local stack

Variables

- Also a set of shared variables, e.g., static variables, shared common blocks, or global heap
 - ❑ Threads communicate implicitly by writing and reading shared variables
 - ❑ Threads **coordinate** by synchronizing on shared variables





Amdahl's Law

- Describes the upper bound of parallel speedup (scaling)
- Helps think about the effects of overhead

Gene M. Amdahl, “*Validity of the Single-Processor Approach to Achieving Large Scale Computing Capabilities*”, 1967

Amdahl's law (Amdahl's speedup model)

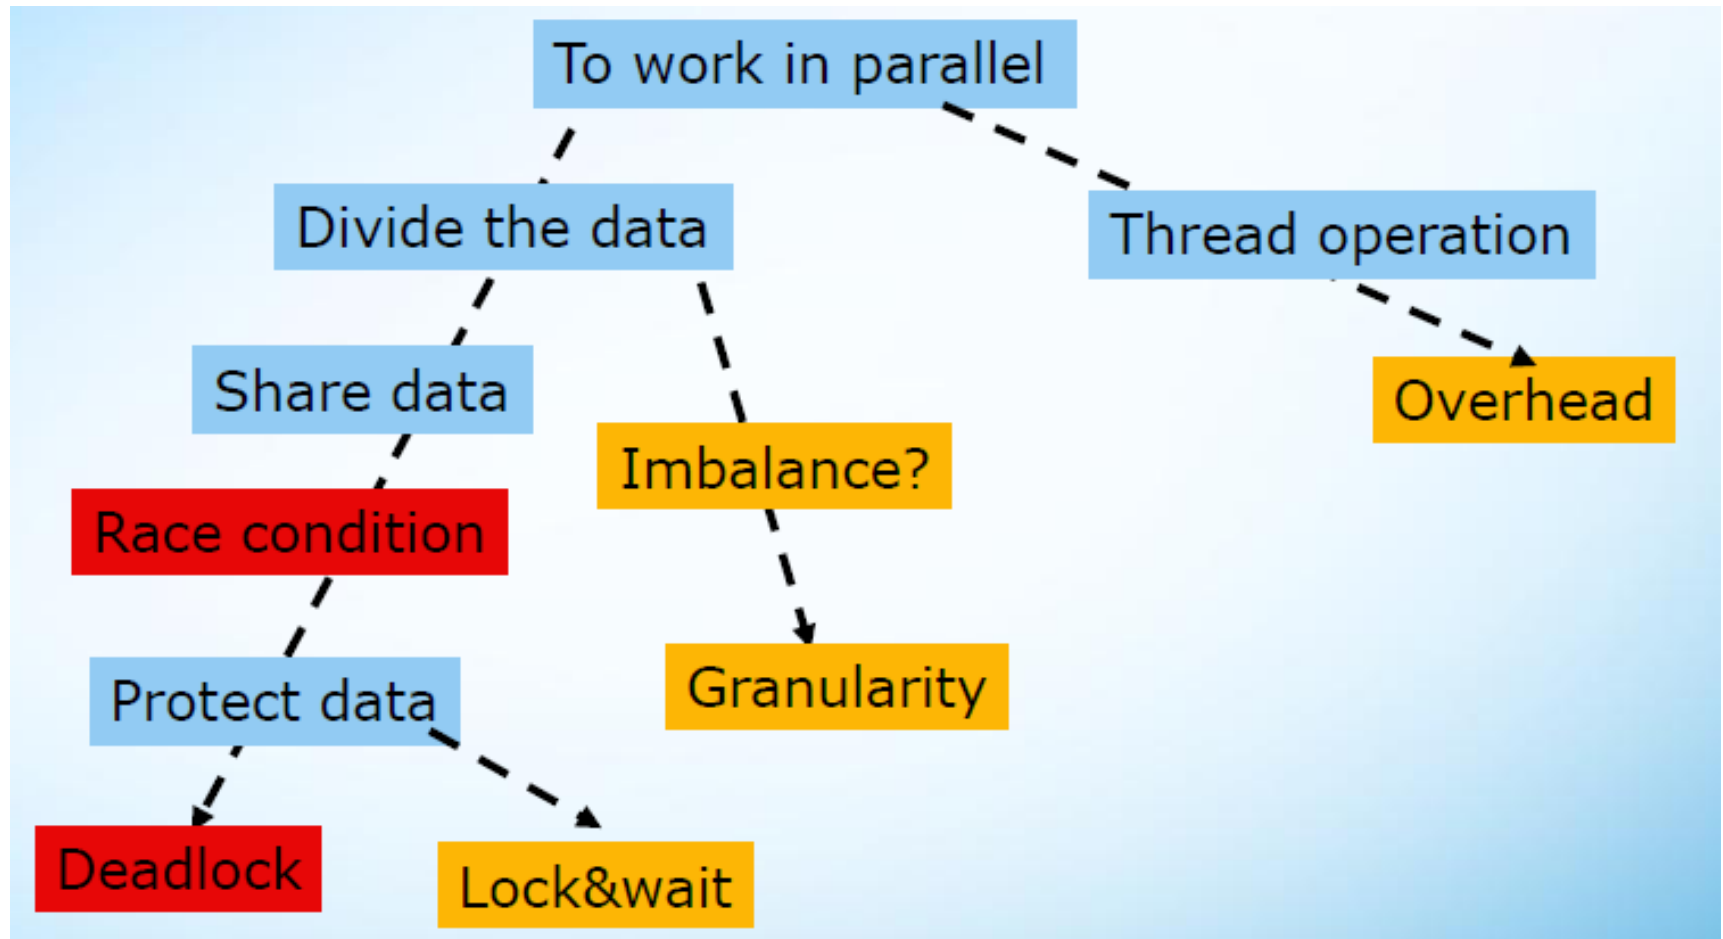
$$Speedup_{Amdahl} = \frac{1}{(1-f) + \frac{f}{n}}$$

$$\lim_{n \rightarrow \infty} Speedup_{Amdahl} = \frac{1}{1-f}$$

f is the parallel portion

Implications


Where Are the Problems From?



Remove the error

Tune for high speedup

Processes and Threads



- Modern operating systems load programs as processes
 - Resource holder
 - Execution
- A process starts executing at its entry point as a thread
- Threads can create other threads within the process
- All threads within a process share code & data segments

Decomposition

➤ Data decomposition

- ❑ Break the entire dataset into smaller, discrete portions, then process them in parallel
- ❑ Folks eat up a cake

➤ Task decomposition

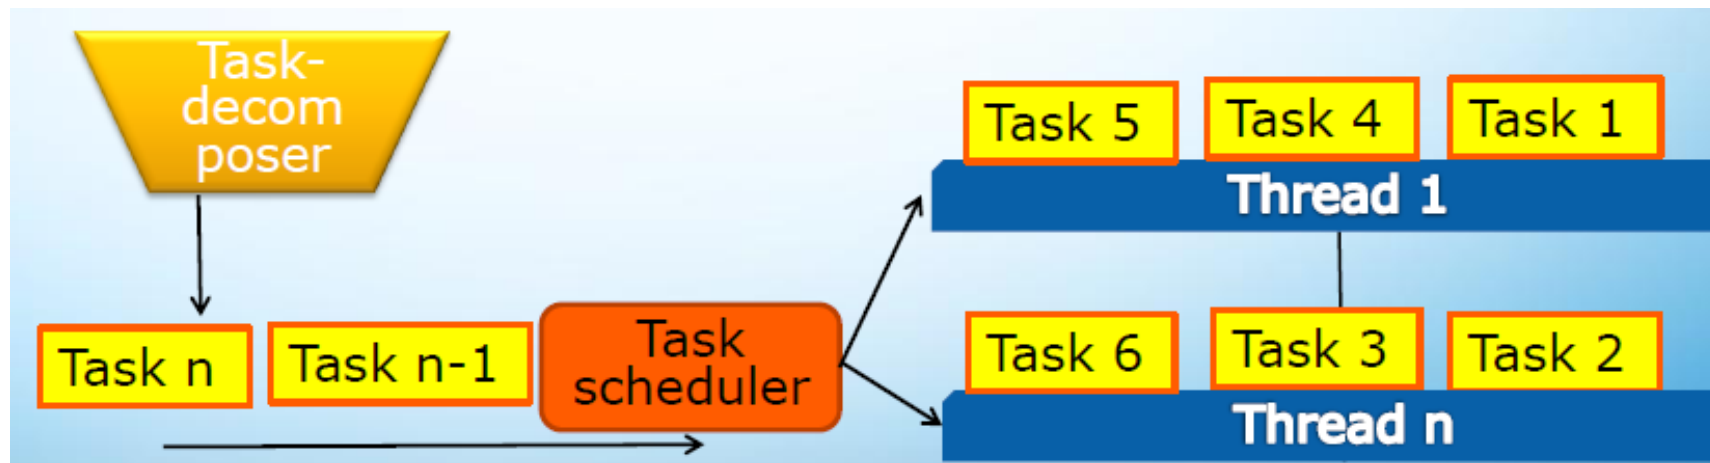
- ❑ Divide the whole task based on natural set of **independent** sub-tasks
- ❑ Folks play a symphony

➤ Considerations

- ❑ Cause less or no share data
- ❑ Avoid the dependency among sub-tasks, otherwise become pipeline

Task and Thread

- A task consists the data and its process, and task scheduler will attach it to a thread to be executed
- Task operation is much cheaper than threading operation
- Ease to balance workload among threads by stealing
- Suit for list, tree, map data structure





Task and Thread

➤ Considerations

- ❑ Many more tasks than threads
 - More flexible to schedule the task
 - Easy to balance workload
- ❑ Amount of computation within a task must be large enough to offset overhead of managing task and thread
- ❑ Static scheduling
 - Tasks are collections of separate, independent function calls or are loop iterations
- ❑ Dynamic scheduling
 - Task execution length is variable and is unpredictable
 - May need an additional thread to manage a shared structure to hold all tasks

Race Conditions

- Threads “*race*” against each other for resources
 - ❑ Execution order is assumed but cannot be guaranteed
- Storage conflict is most common
 - ❑ Concurrent access of same memory location by multiple threads, at least one thread is writing
- Determinacy race and data race
- May not be apparent at all times
- Considerations
 - ❑ Control shared access with critical regions
 - Mutual exclusion and synchronization, critical session, atomic
 - ❑ Scope variables to be local to threads
 - Have a local copy for shared data
 - Allocate variables on thread stack

Deadlock

- 2 or more threads wait for each other to release a resource
- A thread waits for a event that never happen, like suspended lock
- Most common cause is locking hierarchies
- Considerations
 - ❑ Always lock and un-lock in the same order, and avoid hierarchies if possible
 - ❑ Use atomic

```
DWORD WINAPI threadA(LPVOID arg)
{
    EnterCriticalSection(&L1);
    EnterCriticalSection
    processA(data1, da
    LeaveCriticalSection(&L2);
    LeaveCriticalSection(&L1);
    return(0);
}
```

ThreadB: L2, then L1

ThreadA: L1, then L2

```
DWORD WINAPI threadB(LPVOID arg)
{
    EnterCriticalSection(&L2);
    EnterCriticalSection(&L1);
    processB(data2, data1);
    LeaveCriticalSection(&L1);
    LeaveCriticalSection(&L2);
    return(0);
}
```


Thread Safe Routine/Library

- It functions correctly during simultaneous execution by multiple threads
- Non-thread-safe indicators
 - ❑ Access global/static variables or the heap
 - ❑ Allocate/reallocate/free resources that have global scope (files)
 - ❑ Indirect accesses through handles and pointers
- Considerations
 - ❑ Any variables changed must be local to each thread
 - ❑ Routines can use mutual exclusion to avoid conflicts with other threads

**It is better to make a routine reentrant
than to add synchronization
Avoids potential overhead**

Imbalanced Workload

- All threads process the data in same way, but one thread is assigned more work, thus require more time to complete it and impact overall performance
- Considerations
 - ▣ Parallelize the inner loop
 - ▣ Incline to fine-grained
 - ▣ Choice the proper algorithm
 - ▣ Divide and conquer, master and worker, work-stealing



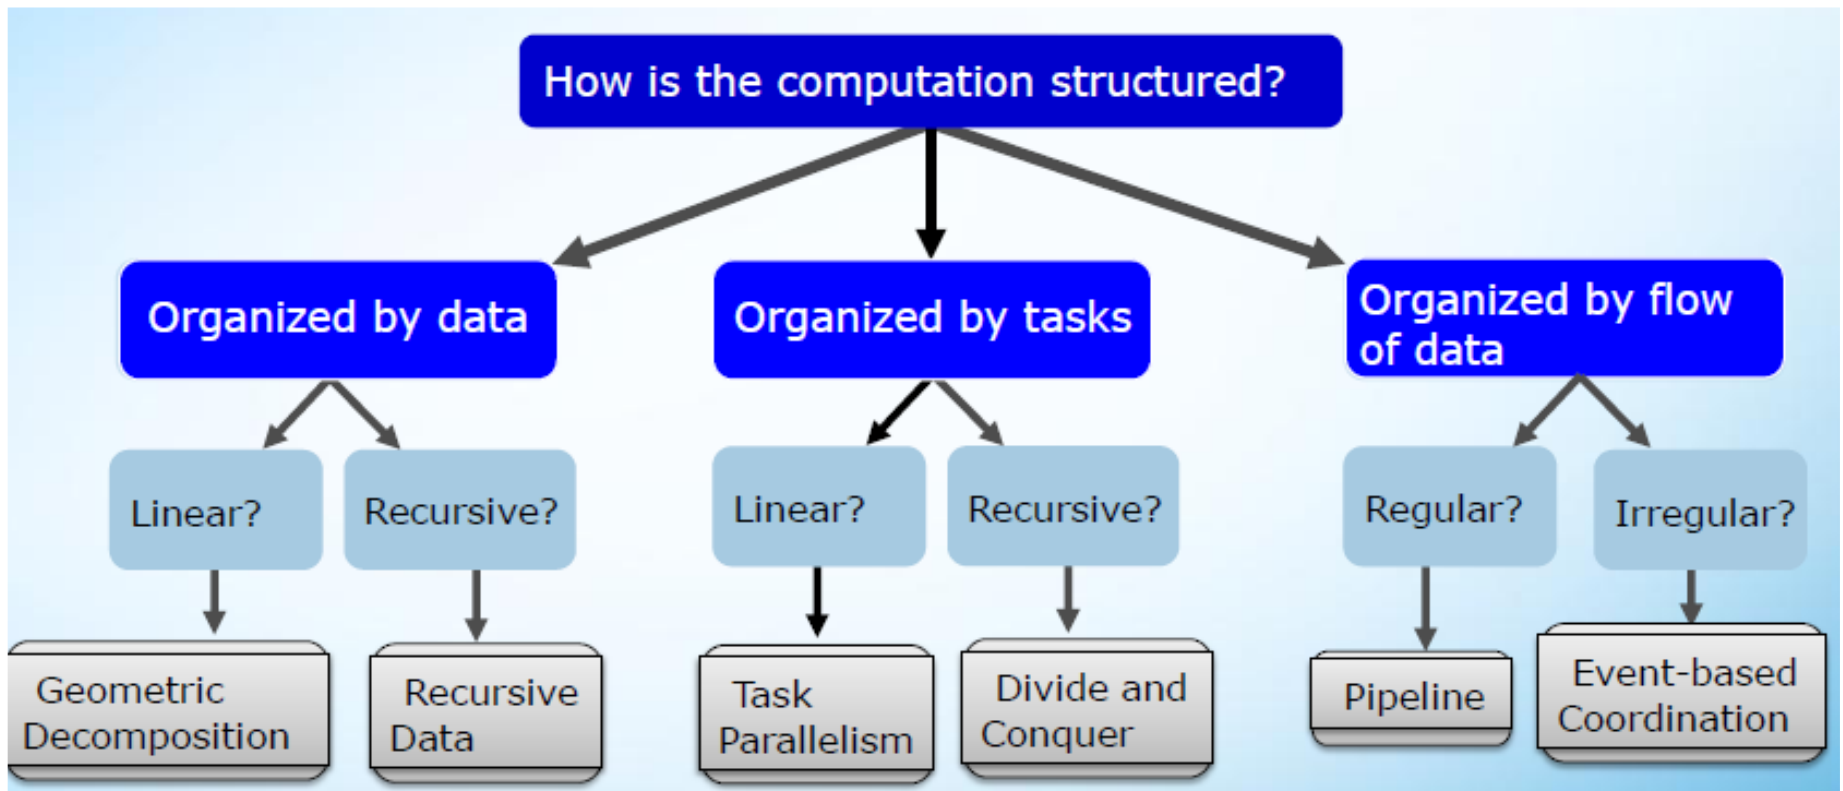
Granularity

- An extent to which a larger entity is subdivided
- Coarse-grained means fewer and larger components
- Fine-grained means more and smaller components
- **Consideration**
 - ❑ Fine-grained will increase the workload for task scheduler
 - ❑ Coarse-grained may cause the workload imbalance
 - ❑ Benchmark to set the proper granularity

Lock & Wait

- Protect shared data and ensure tasks executed in right order
- Improper usage causes a side-effect
- Considerations
 - ❑ Choose appropriate synchronization primitives
 - `tbb::atomic`, `InterlockedIncrement`, `EnterCriticalSection`...
 - ❑ Use non-blocking locks
 - `TryEnterCriticalSection`, `pthread_mutex_try_lock`
 - ❑ Reduce lock granularity
 - ❑ Don't be a lock hub
 - ❑ Introduce a concurrent container for shared data

Parallel Algorithm



A Generic Development Cycle (1)

- Analysis
 - ❑ Find the hotspot and understand its logic
- Design
 - ❑ Identify the concurrent tasks and their dependencies
 - ❑ Decompose the whole dataset with minimal overhead of sharing or data movement between tasks
 - ❑ Introduce the proper parallel algorithm
 - ❑ Use proved parallel implementations
 - ❑ Memory management
 - Avoid heap contention among threads
 - Use thread-local storage to reduce synchronization
 - Detecting memory saturation in threaded applications
 - Avoid and identifying false sharing among threads

A Generic Development Cycle (2)

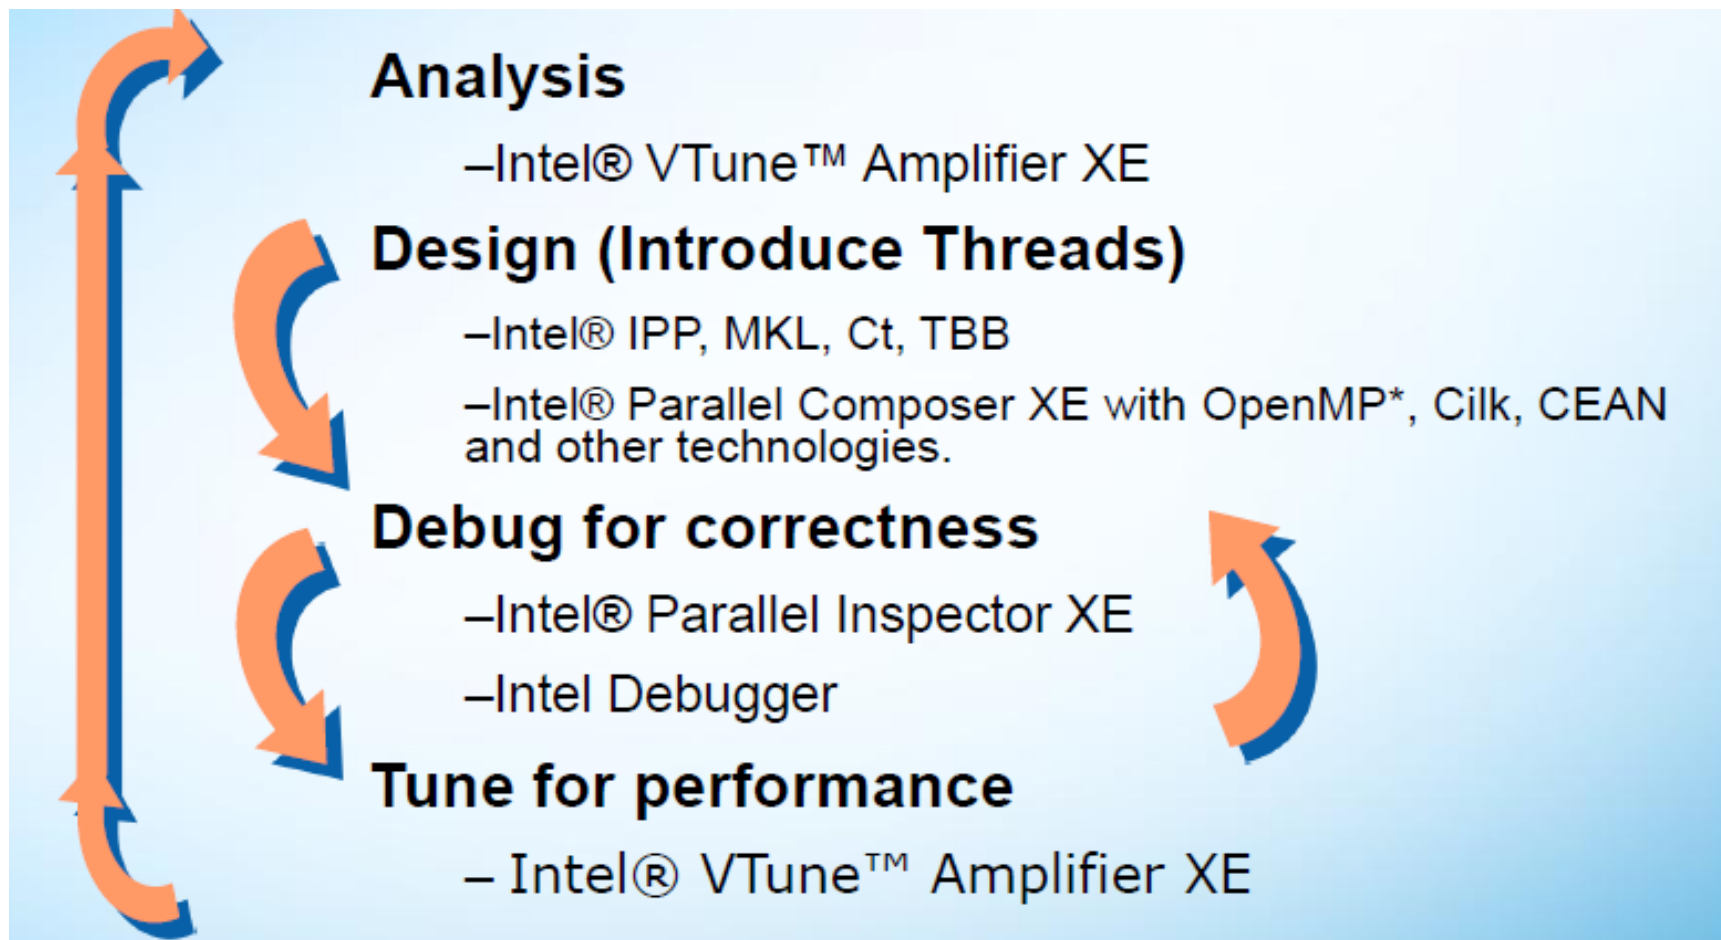
➤ **Debug for correctness**

- ❑ Detect race conditions, deadlock, & memory issues (LLVM , Soot)

➤ **Tune for performance**

- ❑ Balance the workload
- ❑ Adjust lock & wait
- ❑ Reduce thread operation overhead
- ❑ Set the right granularity
- ❑ Benchmark for scalability

Intel Generic Development Cycle





Summary

- **Threading applications require multiple iterations of designing, debugging, and performance tuning steps**
- **Use tools to improve productivity**
- **Unleash the power of dual-core and multi-core processors**




Parallel programming models

MESSAGE PASSING MODEL

Message Passing Architectures


- Complete computer as building block, including I/O
 - ❑ Communication via explicit I/O operations
- Programming model
 - ❑ **directly access** only **private address space** (local memory)
 - ❑ **communicate** via explicit messages (**send/receive**)
- High-level block diagram similar to distributed-mem SAS
 - ❑ But communication integrated at IO level, need not put into memory system
 - ❑ Easier to build than scalable SAS
- Programming model further from basic hardware ops
 - ❑ Library or OS intervention

Message Passing Abstraction



- Send specifies buffer to be transmitted and sending process
- Recv specifies receiving process and application storage to receive into
- Memory to memory copy, but need to name processes
- Optional tag on send and matching rule on receive
- Many overheads: copying, buffer management, protection


Message Passing Abstraction



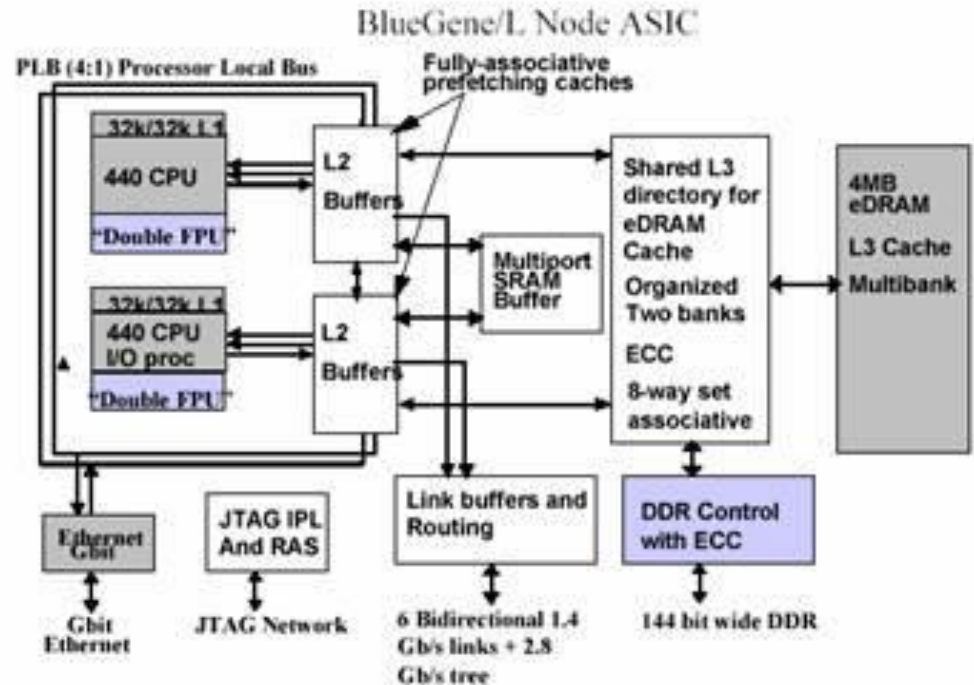
- Send specifies buffer to be transmitted and sending process
- Recv specifies receiving process and application storage to receive into
- Memory to memory copy, but need to name processes
- Optional tag on send and matching rule on receive
- Many overheads: copying, buffer management, protection

Evolution of Message Passing

- **Early machines: FIFO on each link**
 - ❑ Hardware close to programming model
 - synchronous ops
 - ❑ Replaced by DMA, enabling non-blocking ops
 - Buffered by system at destination until recv
- **Diminishing role of topology**
 - ❑ Store & forward routing: topology important
 - ❑ Introduction of pipelined routing made it less so important
 - ❑ Cost is in node network interface
 - ❑ Simplifies programming



Example: IBM Blue Gene/L



Nodes: 2 PowerPC 440s; everything except DRAM on one chip



Toward Architectural Convergence

- Evolution and role of software have blurred boundary
 - ❑ Send/recv supported on SAS machines via buffers
 - ❑ Can construct global address space on MP using hashing
 - ❑ Page-based (or fine-grained) shared virtual memory

- Programming models distinct, but organizations converging
 - ❑ Nodes connected by general network and communication assists
 - ❑ Implementations also converging, at least in high-end machines

Implementations

- From a programming perspective
 - ❑ Message passing implementations usually comprise a library of subroutines
 - ❑ Calls to these subroutines are imbedded in source code
 - ❑ The programmer is responsible for determining all parallelism
- Historically, a variety of message passing libraries have been available since the 1980s. These implementations differed substantially from each other making it difficult for programmers to develop portable applications
- In 1992, the MPI Forum was formed with the primary goal of establishing a standard interface for message passing implementations

Implementations

- Part 1 of the **Message Passing Interface (MPI)** was released in 1994. Part 2 (MPI-2) was released in 1996. Both MPI specifications are available on the web at <http://wwwunix.mcs.anl.gov/mpi/>
- MPI is now the *de facto* industry standard for message passing, replacing virtually all other message passing implementations used for production work
- MPI implementations exist virtually for all popular parallel computing platforms. Not all implementations include everything in both MPI-1 and MPI-2

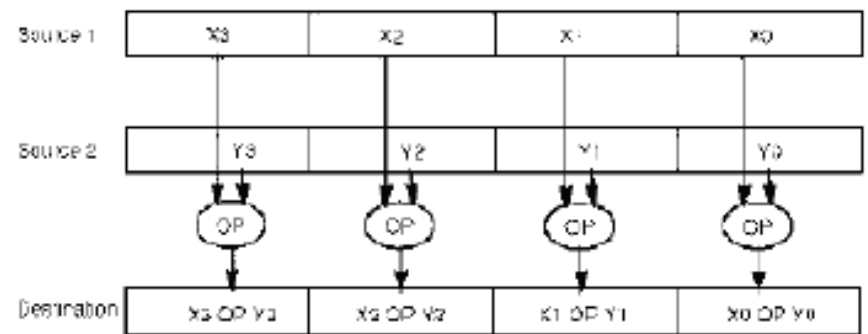



Parallel programming models

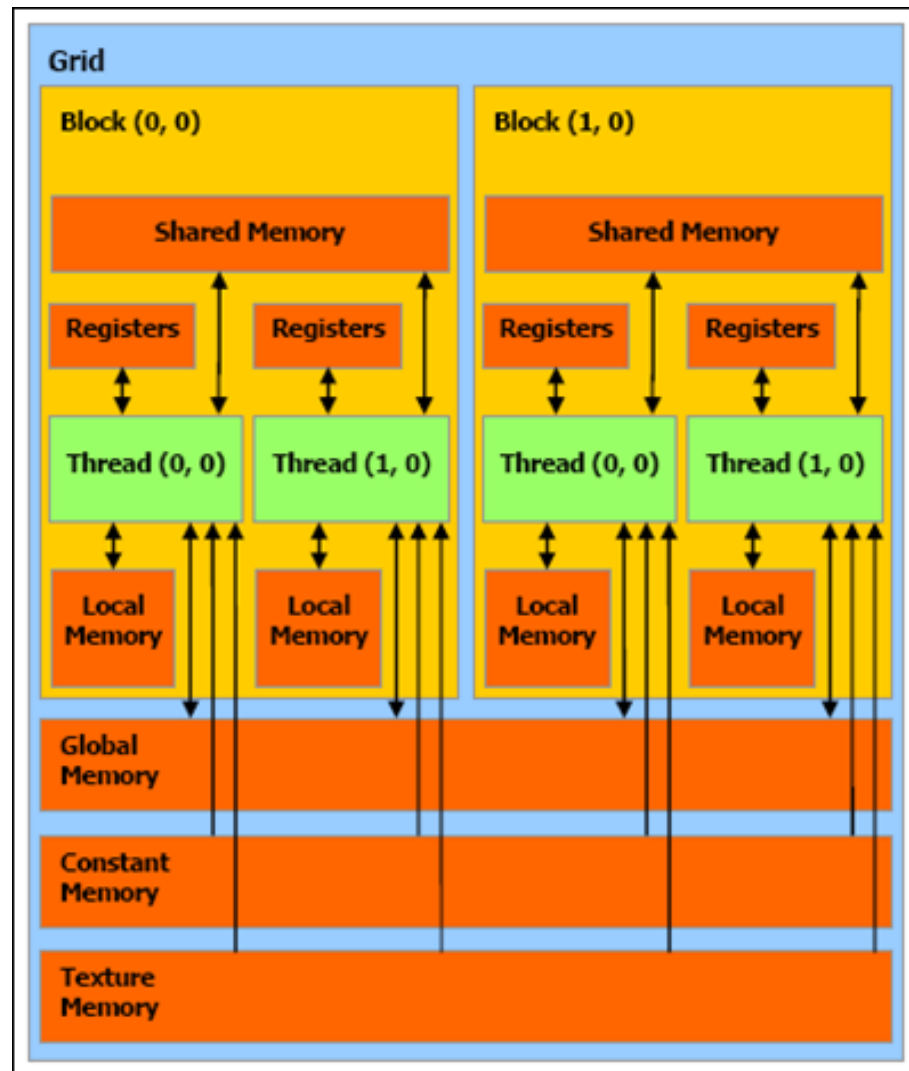
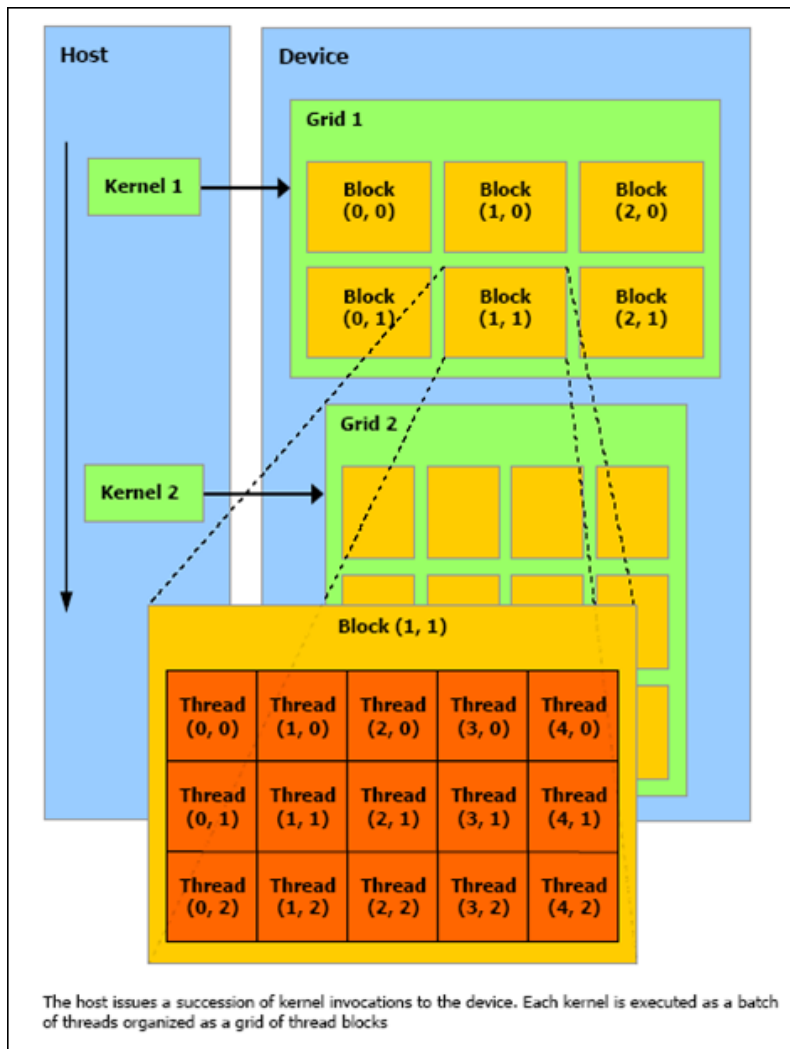
GPGPU PROGRAMMING MODEL

CUDA Goals: SIMD Programming

- Hardware architects love SIMD, since it permits a very space and energy-efficient implementation
- However, standard SIMD instructions on CPUs are inflexible, and difficult to use, difficult for a compiler to target
- CUDA thread abstraction will provide programmability at the cost of additional hardware



CUDA Programming Model






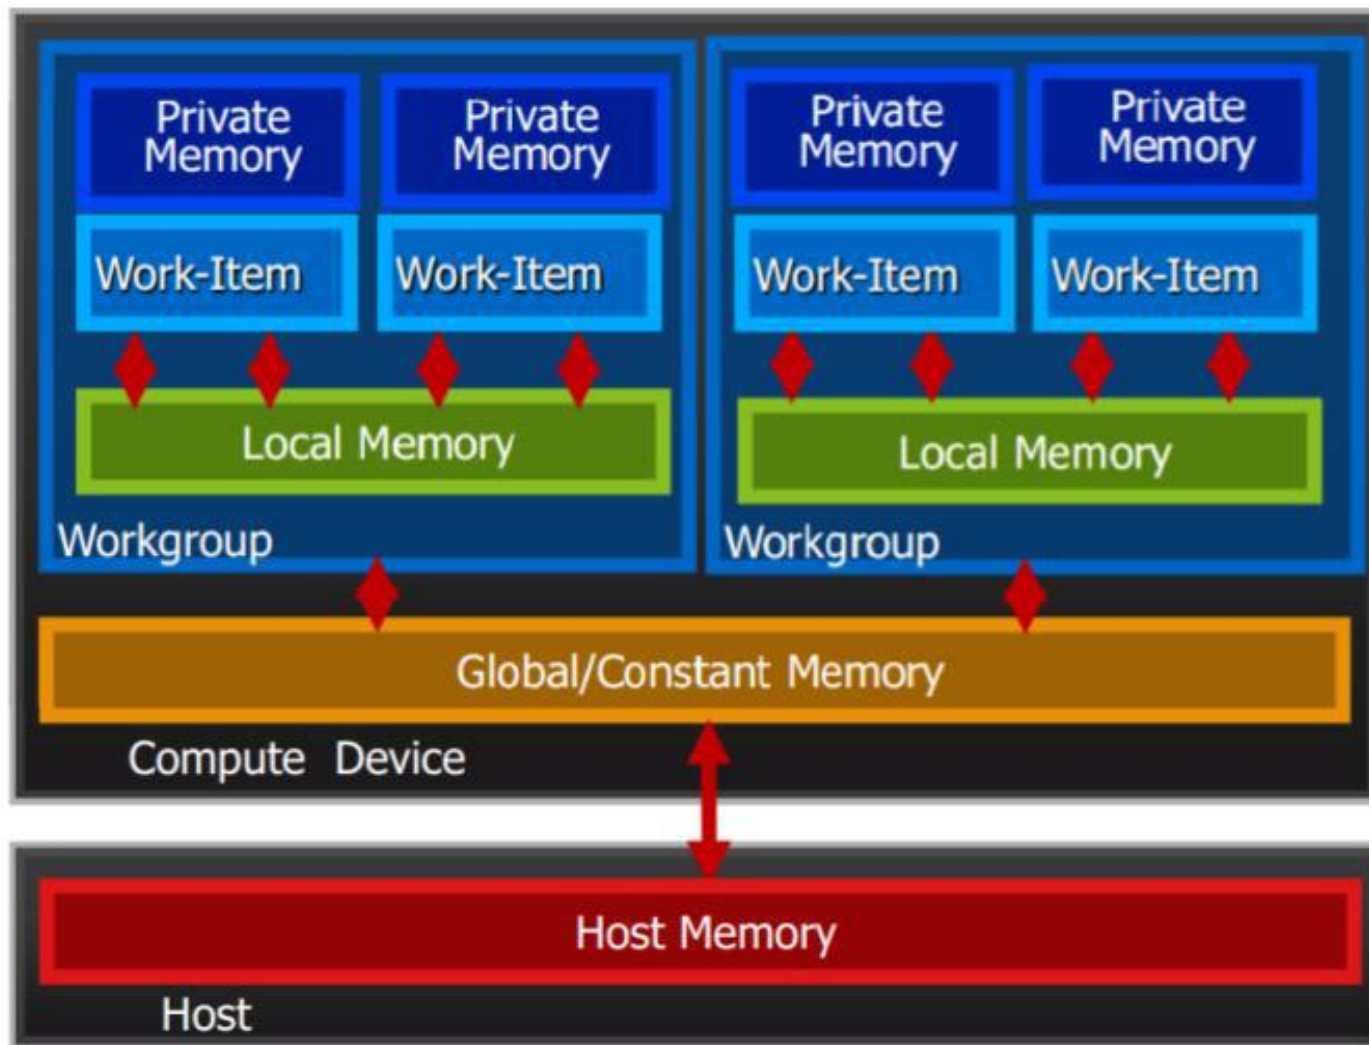
OpenCL Programming Model

- OpenCL is a framework for writing programs that execute across heterogeneous platforms consisting of CPUs, GPUs, DSPs, FPGAs and other processors or hardware accelerators
- Data Parallel - SPMD
 - ❑ Work-items in a work-group run the same program
 - ❑ Update data structures in parallel using the work-item ID to select data and guide execution
- Task Parallel
 - ❑ One work-item per work group ... for coarse grained task-level parallelism
 - ❑ Native function interface: trap-door to run arbitrary code from an OpenCL command-queue

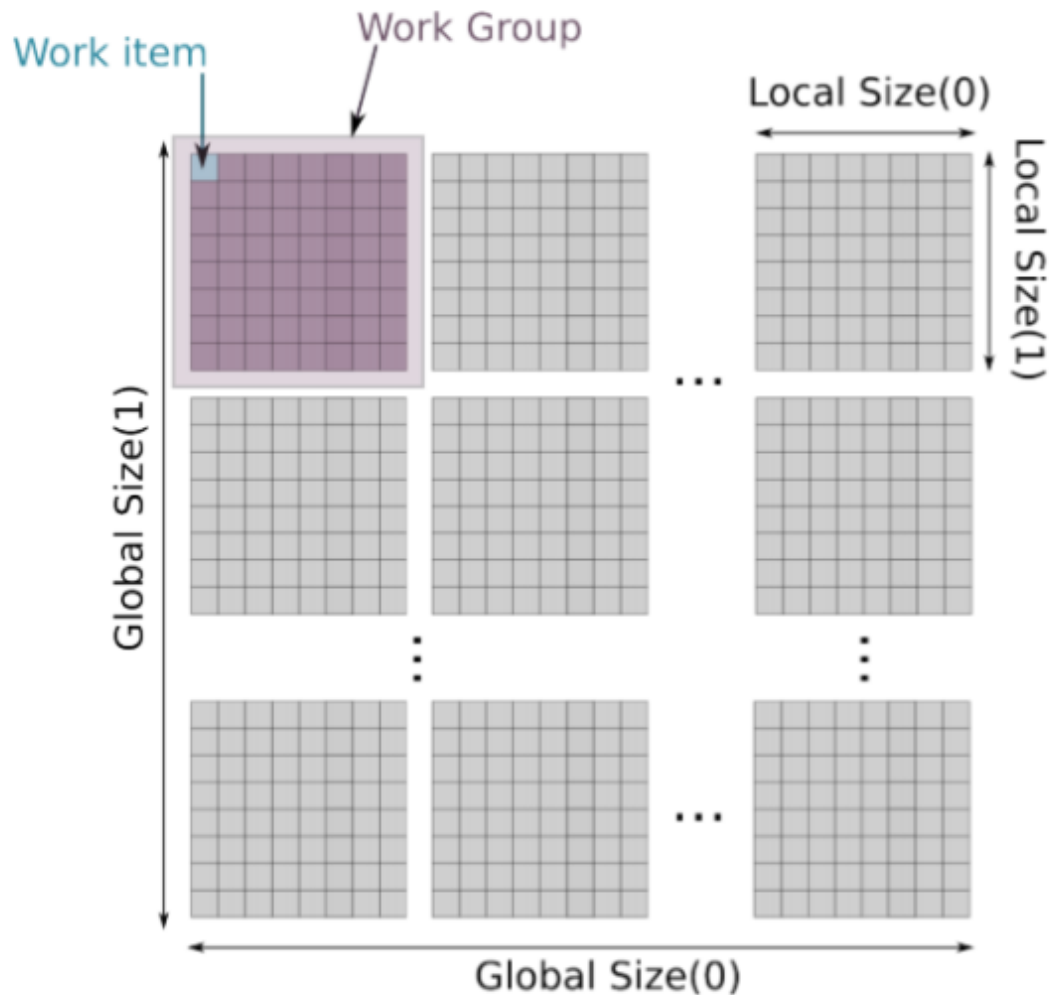
OpenCL Platform Model




OpenCL Memory Model



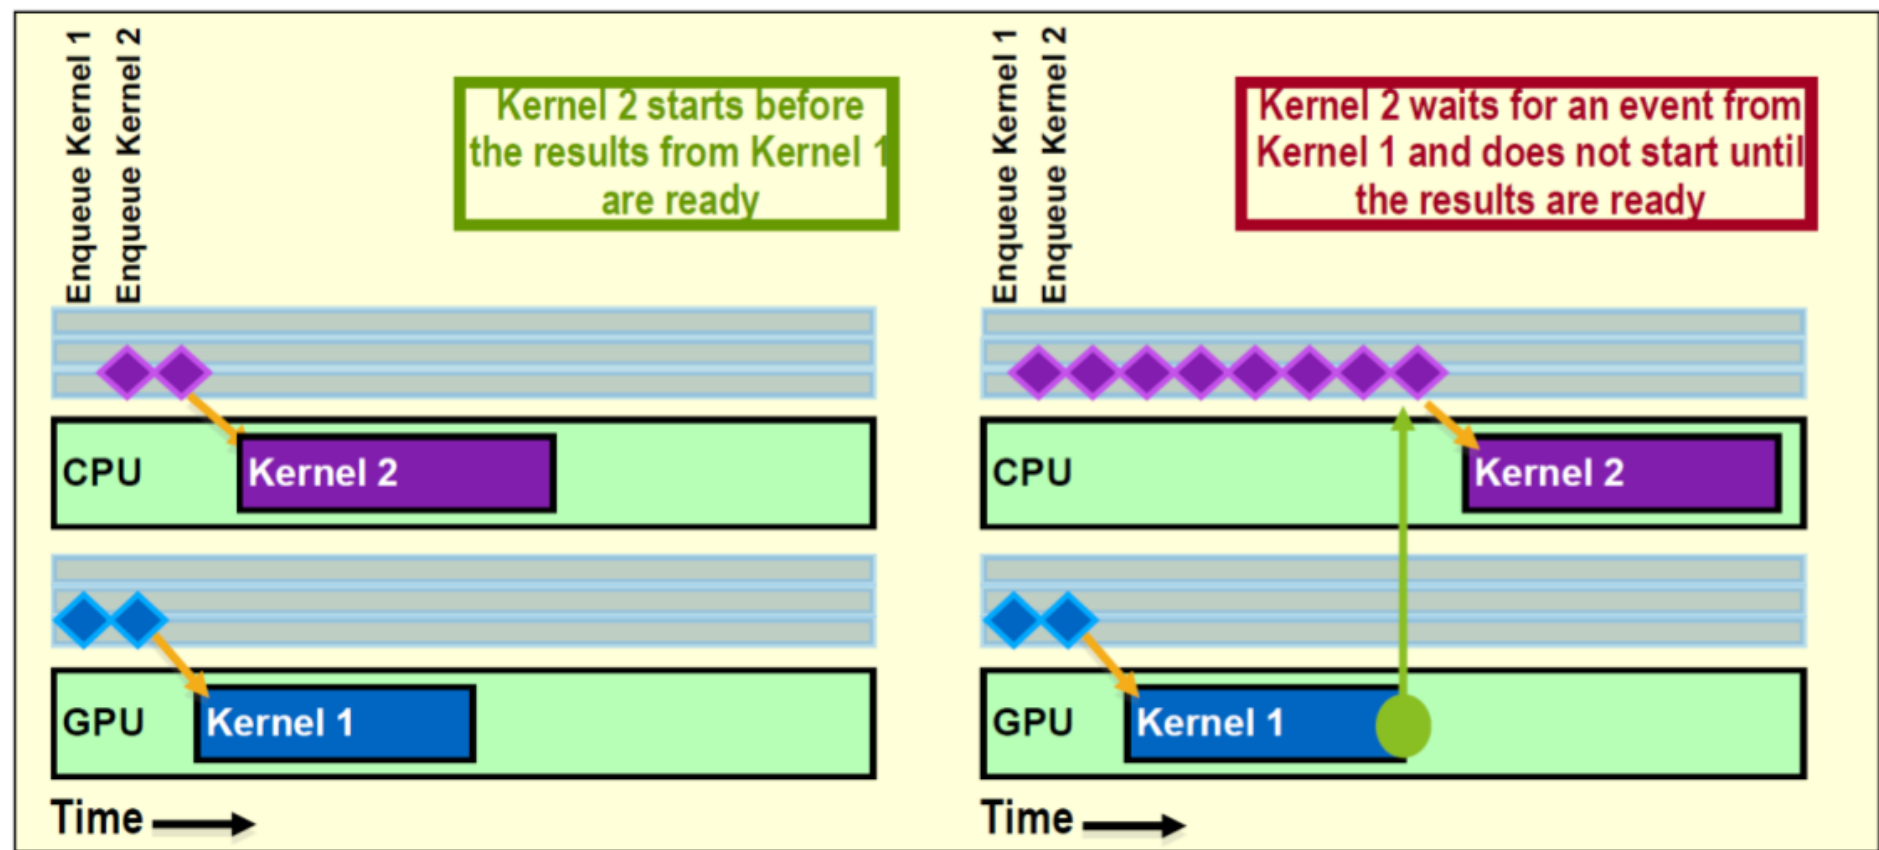
2D Data-Parallel Execution in OpenCL



OpenCL Work-group / Work-unit Structure



Concurrency Control with OpenCL Event-Queueing



*Functions executed on an OpenCL device are called kernels



OpenCL's Two Styles of Data Parallelism

- Explicit SIMD data parallelism
 - ❑ The kernel defines one stream of instructions
 - ❑ Parallelism from using wide vector types
 - ❑ Size vector types to match native HW width
 - ❑ Combine with task parallelism to exploit multiple cores
- Implicit SIMD data parallelism (i.e. shader-style)
 - ❑ Write the kernel as a “scalar program”
 - ❑ Use vector data types sized naturally to the algorithm
 - ❑ Kernel automatically mapped to SIMD-compute-resources and cores by the compiler/runtime/hardware

Both approaches are viable CPU options



Parallel programming models

Data Parallel Systems



Data Parallel Systems

➤ Programming model

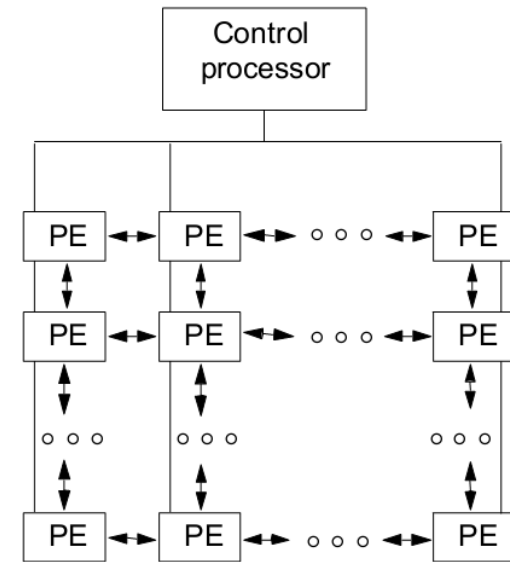
- ❑ Operations performed in parallel on each element of data structure
- ❑ Logically single thread of control, performs sequential or parallel steps
- ❑ Conceptually, a processor associated with each data element

➤ Architectural model

- ❑ Array of many simple, cheap processors with little memory each
 - Processors don't sequence through instructions
- ❑ Attached to a control processor that issues instructions
- ❑ Specialized and general communication, cheap global synchronization

➤ Original motivation

- ❑ Matches simple differential equation solvers
- ❑ Centralize high cost of instruction fetch & sequencing



Application of Data Parallelism

➤ Example

- ❑ Each PE contains an employee record with his/her salary

If salary > 100K then

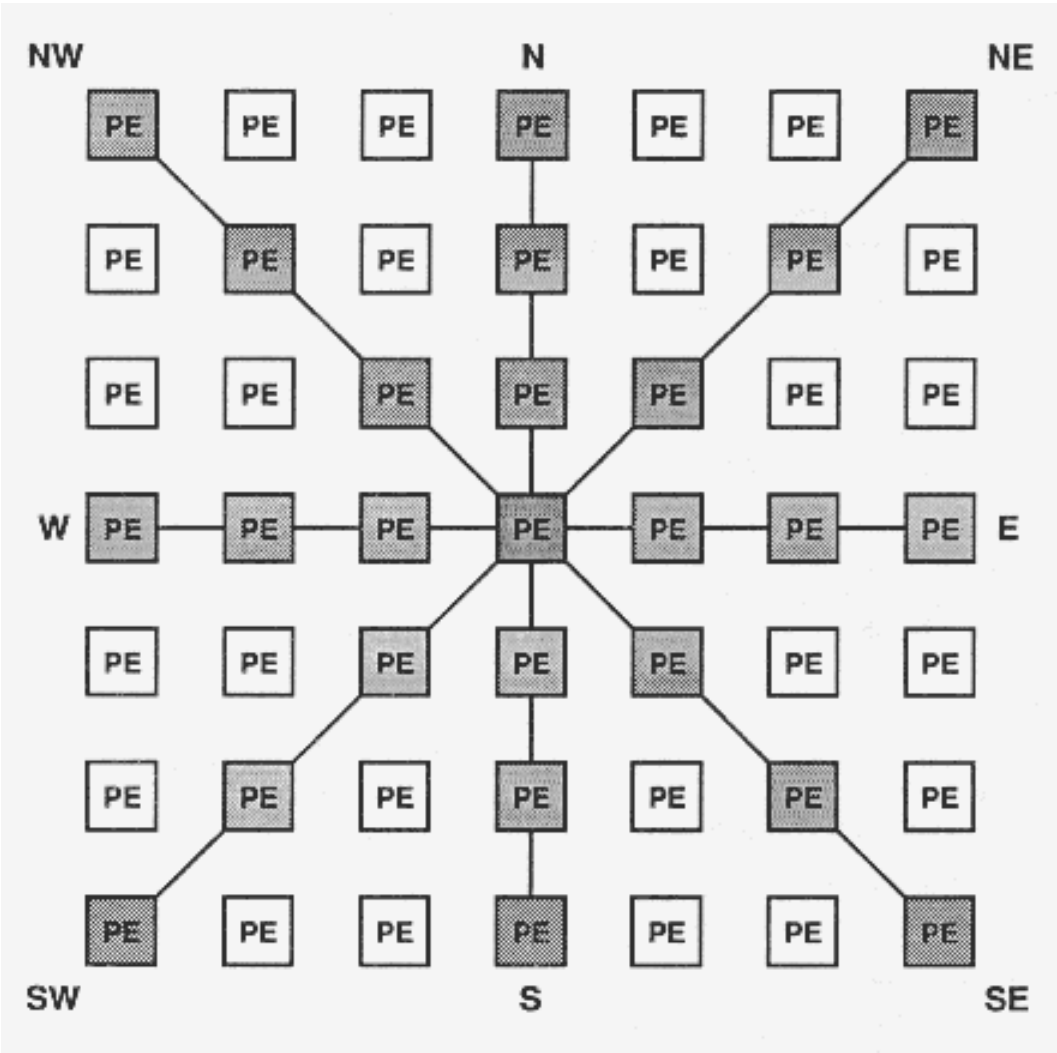

salary = salary * 1.05

else

salary = salary * 1.10

- ❑ Logically, the whole operation is a single step
 - ❑ Some processors enabled for arithmetic operation, others disabled
- ## ➤ Other examples
- ❑ Finite differences, linear algebra, ...
 - ❑ Document searching, graphics, image processing, ...
- ## ➤ Example machines
- ❑ Thinking Machines CM-1, CM-2 (and CM-5)
 - ❑ Maspar MP-1 and MP-2

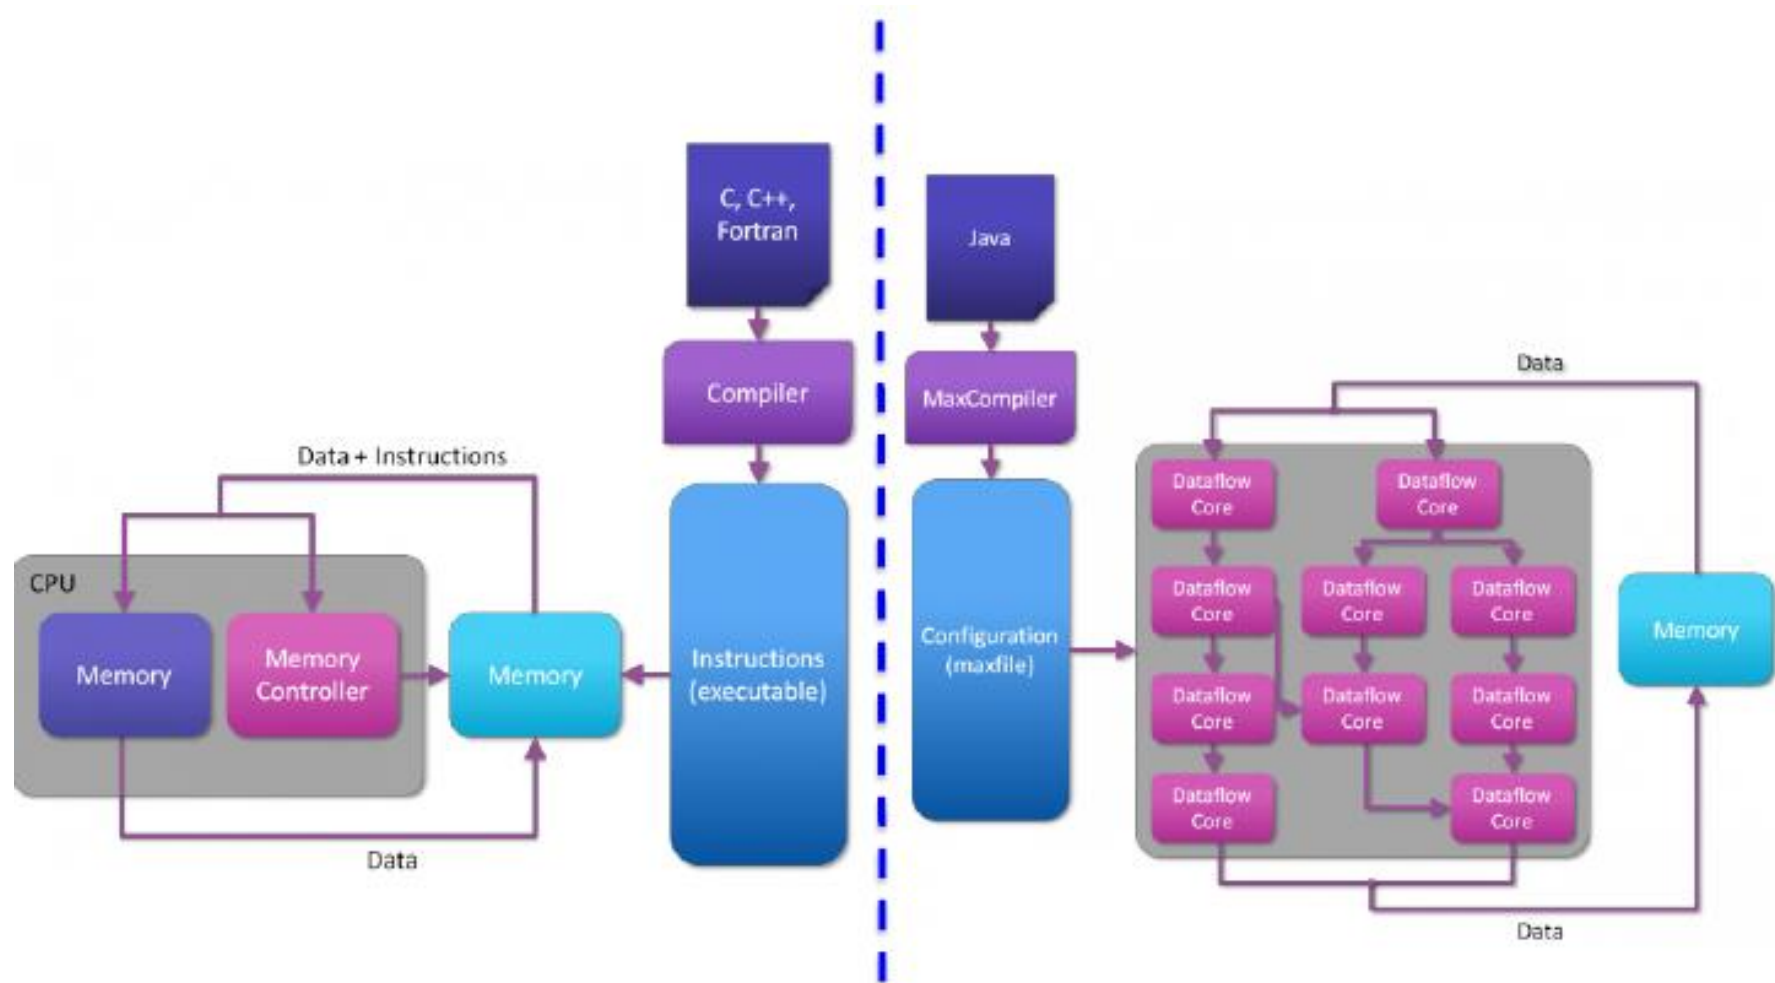
Maspar MP Architecture



Dataflow Architecture



- Non-von Neumann models of computation, architecture, and languages
- Programs are not attached to a program counter
- Executability and execution of instructions is solely determined based on the availability of input arguments to the instructions
- Order of instruction execution is unpredictable: i. e. behavior is indeterministic
- Static and Dynamic dataflow machines
 - ▣ Static dataflow machines: use conventional memory addresses as data dependency tags
 - ▣ Dynamic dataflow machines: use content-addressable memory (CAM)

Computing with Control Flow/ Data Flow Cores



Control Flow vs. Data Flow

```
a := x + y
b := a × a
c := 4 - a
```



Evolution and Convergence

- Rigid control structure (SIMD in Flynn taxonomy)
 - ❑ SISD = uniprocessor , MIMD= multiprocessor
- Popular when cost savings of centralized sequencer high
 - ❑ 60s when CPU was a cabinet; replaced by vectors in mid-70s
 - ❑ Revived in mid-80s when 32-bit data path slices just fit on chip
 - ❑ No longer true with modern microprocessors
- Other reasons for demise
 - ❑ Simple, regular applications have good locality, can do well anyway
 - ❑ Loss of applicability due to hardwiring data parallelism
 - MIMD machines as effective for data parallelism and more general
- Programming model converges with SPMD (single program multiple data)
 - ❑ Contributes need for fast global synchronization
 - ❑ Structured global address space, implemented with either SAS or MP

1. Why is it difficult to construct a true shared-memory computer? What is the minimum number of switches for connecting p processors to a shared memory with b words (where each word can be accessed independently)?

2. Consider a memory system with a level 1 cache of 32 KB and DRAM of 512 MB with the processor operating at 1 GHz. The latency to L1 cache is one cycle and the latency to DRAM is 100 cycles. In each memory cycle, the processor fetches four words (cache line size is four words). What is the peak achievable performance of a dot product of two vectors? Note: Where necessary, assume an optimal cache placement policy.

```
/* dot product loop */  
for (i = 0; i < dim; i++)  
    dot_prod += a[i] * b[i];
```




❖ Project

Implement a multi-access threaded queue with multiple threads inserting and multiple threads extracting from the queue. Use mutex-locks to synchronize access to the queue. Document the time for 1000 insertions and 1000 extractions each by 64 insertion threads (producers) and 64 extraction threads (consumers).



Thank You !