



# 并行与分布式计算

Parallel & Distributed Computing

陈鹏飞  
数据科学与计算机学院  
2018-04-13



# Lecture 5 — OpenMP Programming

Pengfei Chen

School of Data and Computer Science

April 13, 2018

## ***Outline:***

**1**

### **OpenMP Overview**

**2**

### **Correctness**

**3**

### **Task**

**4**


### **Performance**



**OpenMP programming**


# Overview

# Architecture for Shared Memory Model



**Uniform Memory Access**

**Non-uniform Memory Access**



# Thread Base Parallelism

- OpenMP programs accomplish parallelism (显式) exclusively (仅仅) through the use of threads
- A thread of execution is the smallest unit of processing that can be scheduled by an operating system
  - ❑ The idea of a subroutine that can be scheduled to run autonomously might help explain what a thread is
- Threads exist within the resources of a single process
  - ❑ Without the process, they cease (停止) to exist
- Typically, the number of threads match the number of machine processors/cores
  - ❑ However, the actual use of threads is up to the application



# *Explicit Parallelism*

- OpenMP is an explicit (not automatic) programming model, offering the programmer full control over parallelization
- Parallelization can be as simple as taking a serial program and inserting compiler directives....
- Or as complex as inserting subroutines to set multiple levels of parallelism, locks and even nested locks



# What is OpenMP?

## ➤ An abbreviation for

### □ Short version

- **Open Multi-Processing**（开放多处理过程）

### □ Long version

- **Open** specifications for **Multi-Processing** via collaborative work between interested parties from the hardware and software industry, government and academia



# OpenMP Overview

C\$OMP FLUSH

#pragma omp critical

C\$OMP THREADPRIVATE (/ABC/)

CALL OMP SET NUM THREADS (10)

## *OpenMP: An API for Writing Multithreaded Applications*

- A set of compiler directives and library routines for parallel application programmers
- Greatly simplifies writing multi-threaded (MT) programs in Fortran, C and C++
- Standardizes last 20 years of SMP practice

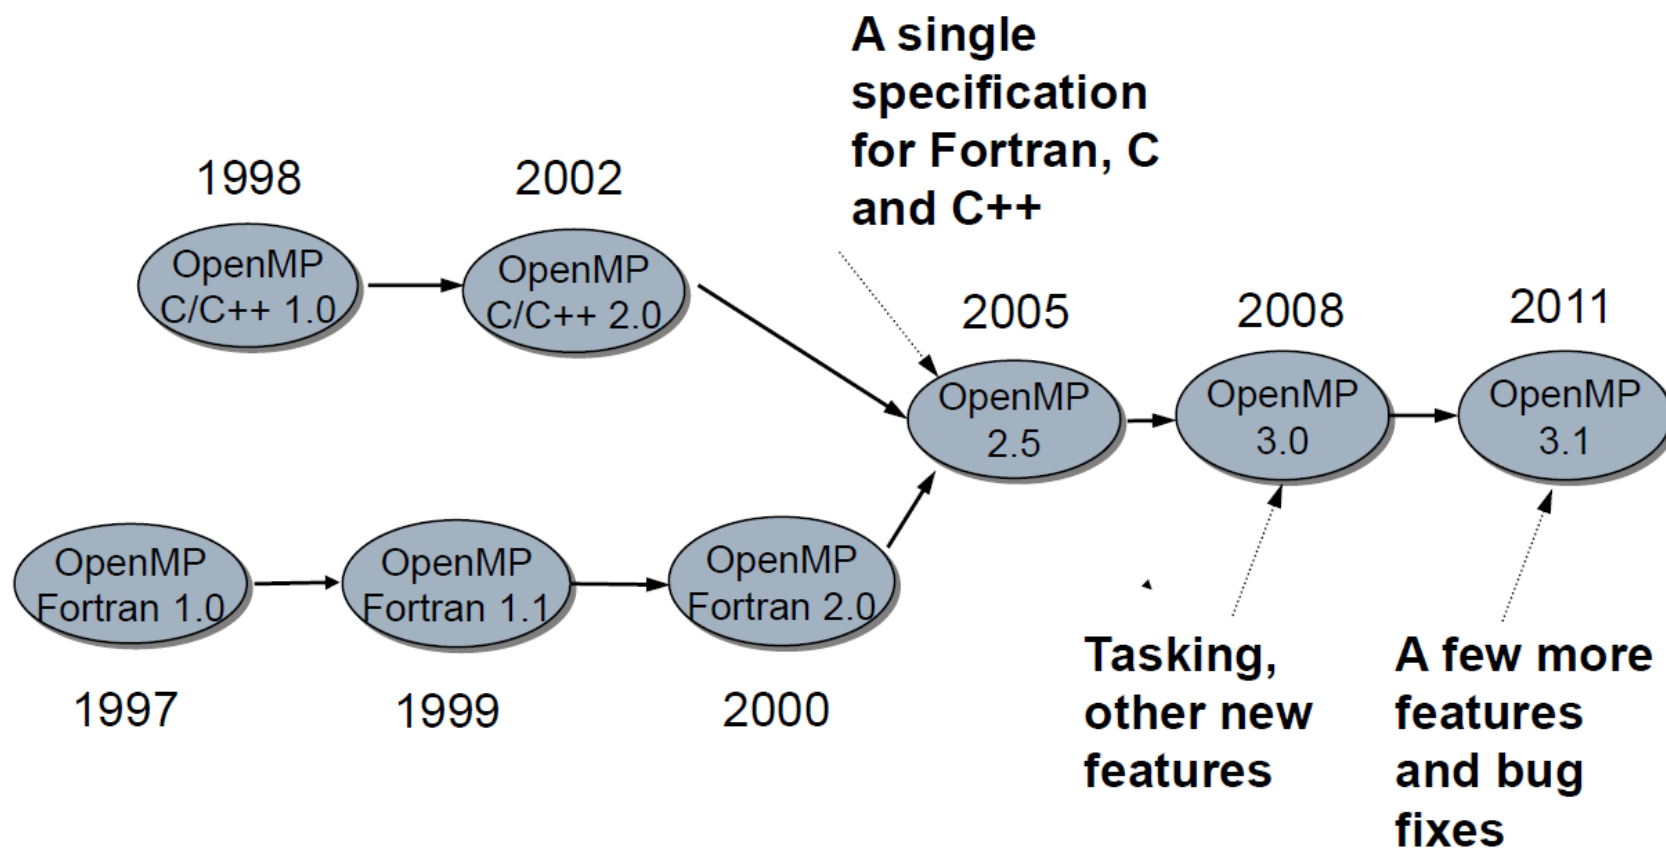
C\$OMP PARALLEL COPYIN (/blk/)

C\$OMP DO lastprivate (XX)

Nthrds = OMP\_GET\_NUM\_PROCS ()

omp\_set\_lock (lck)

# OpenMP Release History



# ***OpenMP Overview: How do Threads Interact?***

- OpenMP is a multi-threading, shared address model
  - ▣ Threads communicate by sharing variables
- Unintended sharing of data causes race conditions
  - ▣ Race condition: when the program's outcome changes as the threads are scheduled differently
- To control race conditions
  - ▣ Use synchronization to protect data conflicts
- Synchronization is expensive
  - ▣ Change how data is accessed to minimize the need for synchronization


# Fork-Join Model

➤ FORK


- ❑ The master thread then creates (or awakens) a team of parallel threads.

➤ JOIN

- ❑ When the team threads complete the statements in the parallel region construct, they synchronize and terminate, leaving only the master thread.



# Relating Fork/Join to Code





# ***OpenMP Components***

- **Three API components**
  - ❑ **Compiler directives**
  - ❑ **Runtime library routines**
  - ❑ **Environment variables**



# Syntax of Compiler Directives (指令)

- *pragma*: a C/C++ compiler directive (编译开关)
  - ❑ (other compiler directives: #include, #define, ...)
  - ❑ Stands for “**pragmatic information** (附注信息)”
  - ❑ A way for the **programmer** to communicate with the **compiler**
  - ❑ Pragmas are **handled by the preprocessor**
  - ❑ Compilers are free to **ignore pragmas**
- All OpenMP pragmas have the syntax:
  - ❑ **#pragma omp <directive-name> [clause, ...]**
- Pragmas appear immediately **before** relevant construct





# Hello World

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>
int main (int argc, char *argv[]) {
    int nthreads, tid;
    /* Fork a team of threads giving them their own copies of variables */
    #pragma omp parallel private(nthreads, tid)
    {
        /* Obtain thread number */
        tid = omp_get_thread_num();
        /* Only master thread does this */
        if (tid == 0) {
            nthreads = omp_get_num_threads();
            printf("Number of threads = %d\n", nthreads);
        }
        printf("Hello World from thread = %d\n", tid);
    } /* All threads join master thread and disband */
}
```



## ***Output – Non-deterministic!***

```
[cong@lnxsrvg1 ~/cs133_examples]$ ./a.out
```

```
Hello World from thread = 5
```

```
Hello World from thread = 7
```

```
Number of threads = 16
```

```
Hello World from thread = 0
```

```
Hello World from thread = 8
```

```
Hello World from thread = 3
```

```
Hello World from thread = 1
```

```
Hello World from thread = 6
```

```
Hello World from thread = 15
```

```
Hello World from thread = 14
```

```
Hello World from thread = 10
```

```
Hello World from thread = 2
```

```
Hello World from thread = 9
```

```
Hello World from thread = 11
```

```
Hello World from thread = 13
```

```
Hello World from thread = 12
```

```
Hello World from thread = 4
```

```
[cong@lnxsrvg1 ~/cs133_examples]$ ./a.out
```

```
Hello World from thread = 6
```

```
Hello World from thread = 10
```

```
Hello World from thread = 8
```

```
Number of threads = 16
```

```
Hello World from thread = 0
```

```
Hello World from thread = 5
```

```
Hello World from thread = 14
```

```
Hello World from thread = 15
```

```
Hello World from thread = 13
```

```
Hello World from thread = 12
```

```
Hello World from thread = 4
```

```
Hello World from thread = 3
```

```
Hello World from thread = 9
```

```
Hello World from thread = 2
```

```
Hello World from thread = 7
```

```
Hello World from thread = 1
```

```
Hello World from thread = 11
```



## ***Supplemental: printf & cout***

- C/C++ is not aware of “threads,” but POSIX is.

- ❑ <http://pubs.opengroup.org/onlinepubs/9699919799/functions/flockfile.html>

- ❑ All functions that reference (FILE \*) objects shall behave as if they use flockfile() and funlockfile() internally to obtain ownership of these (FILE \*) objects.

- What about replacing the printf statement ... ?

- ```
printf("Hello World from thread = %d\n", tid);
```

- by**

- ```
cout << "Hello World from thread = " << tid << endl;
```

# OpenMP Compilers Perform the Translations to Threads (e.g. pthreads)

```
int a, b;
main() {
    [ // serial segment
      #pragma omp parallel num_threads(8) private(a) shared(b)
      { [ // parallel segment
        ]
      }
    [ // rest of serial segment
    ]
}
```

Sample OpenMP program

Code inserted by the OpenMP compiler

```
int a, b;
main() {
    [ // serial segment
      for (i = 0; i < 8; i++)
        pthread_create (....., internal_thread_fn_name, ...);
      for (i = 0; i < 8; i++)
        pthread_join (.....);
    ]
    [ // rest of serial segment
    ]
    void *internal_thread_fn_name (void *packaged_argument) {
        int a;
    }
    [ // parallel segment
    ]
}
```

Corresponding Pthreads translation



## Matching Threads with CPUs

- ◆ Function *omp\_get\_num\_procs* returns the number of physical processors available to the parallel program

```
int omp_get_num_procs(void);
```

- ◆ Function *omp\_set\_num\_threads* allow you to set the number of threads that should be active in parallel sections of code

```
void omp_set_num_threads(int t);
```

- The function can be called with different arguments at different points in the program

## *Pragma: parallel for*

### ◆ The compiler directive

`#pragma omp parallel for`

tells the compiler that the *for* loop which immediately follows can be executed in parallel

- The number of loop iterations must be computable at run time before loop executes
- Loop must not contain a *break*, *return*, or *exit*
- Loop must not contain a *goto* to a label outside loop

*Example: parallel for*

```
int a[1000], b[1000], s[1000];
```

```
...
```

```
#pragma omp parallel for
```

```
for (i = 0; i < 1000; i ++)
```

```
    s[i] = a[i] + b[i];
```

- ◆ Threads are assigned an independent set of iterations
- ◆ Threads must wait at the end of construct





## *Which Loop to Make Parallel?*

```
int main() {  
    int i, j, k;  
    float **a, **b;  
    ... // initialize a[][], b[][] as the 1-hop distance matrix  
    for (k = 0; k < N; k++) {  
        for (i = 0; i < N; i++)  
            for (j = 0; j < N; j++)  
                a[i][j] = min(a[i][j], b[i][k] + b[k][j]);  
        ... // copy a[][] to b[][]  
    }  
}
```

## *Which Loop to Make Parallel?*

```
int main() {  
    int i, j, k;  
    float **a, **b;  
    ... // initialize b[][] as the 1-hop distance matrix  
    for (k = 0; k < N; k++)           // Loop-carried dependences  
        for (i = 0; i < N; i++)       // Can execute in parallel  
            for (j = 0; j < N; j++)   // Can execute in parallel  
                a[i][j] = min(a[i][j], b[i][k] + b[k][j]);  
    ... // copy a[][] to b[][]
```

# Minimizing Threading Overhead

- There is a fork/join for every instance of

```
#pragma omp parallel for
```

```
for (...) {
```

```
...
```

```
}
```

- Since fork/join is a source of overhead, we want to maximize the amount of work done for each fork/join; i.e., the *grain size*
- Hence we choose to make the middle loop parallel

# Almost Right, but Not Quite

```
int main() {  
    int i, j, k;  
    float **a , **b;  
    ... // initialize b[][] as the 1-hop distance matrix  
    for (k = 0; k < N; k++) {  
        #pragma omp parallel for  
        for (i = 0; i < N; i++)  
            for (j = 0; j < N; j++)  
                a[i][j] = min(a[i][j], b[i][k] + b[k][j]);  
        ... // copy a[][] to b[][]  
    }
```

Problem: j is a shared variable



# *Clause: private*

## ➤ Clause

□ An optional, additional component to a pragma

➤ Private clause: directs compiler to make one or more variables

**private**

#pragma omp ... private (<variable list>)



## ***Problem Solved with private Clause***

```
int main() {  
    int i, j, k;  
    float **a , **b;  
    ... // initialize b[][] as the 1-hop distance matrix  
    for (k = 0; k < N; k++) {  
        #pragma omp parallel for private (j)  
        for (i = 0; i < N; i++)  
            for (j = 0; j < N; j++)  
                a[i][j] = min(a[i][j], b[i][k] + b[k][j]);  
        ... // copy a[][] to b[][]  
    }  
}
```

Tell compiler to make  
listed variables private



## *Another Example*

```
int i;  
  
float *a, *b, *c, tmp;  
  
...  
  
for (i = 0; i < N; i++) {  
    tmp = a[i] / b[i];  
    c[i] = tmp * tmp;  
}
```

**Loop is perfectly parallelizable except for shared variable *tmp***





# Solution

```
int i;  
  
float *a, *b, *c, tmp;  
  
...  
  
#pragma omp parallel for private (tmp)  
for (i = 0; i < N; i++) {  
    tmp = a[i] / b[i];  
    c[i] = tmp * tmp;  
}
```

## More About Private Variables (私有变量)

- Each thread **has its own copy of the private variables**
- If  $j$  is declared private, then **inside the *for* loop no thread can access the “other”  $j$  (the  $j$  in shared memory)**
- **No** thread can **use a previously defined value of  $j$**
- **No** thread can **assign a new value to the shared  $j$**
- Private variables are **undefined** at loop entry and loop exit,  
reducing execution time



## *Clause: firstprivate*

- The ***firstprivate* clause** tells the compiler that **the private variable should inherit the value of the shared variable upon loop entry**
- The value is assigned once per thread, not once per loop iteration

## *Example: firstprivate*

```
a[0] = 0.0;
```

```
for (i = 1; i < N; i++)
```

```
    a[i] = alpha(i, a[i-1]);
```

```
#pragma omp parallel for firstprivate (a)
```

```
for (i = 0; i < N; i++)
```

```
    b[i] = beta(i, a[i]);
```

```
    a[i] = gamma(i);
```

```
    c[i] = delta(a[i], b[i]);
```

```
}
```

## Clause: *firstprivate*

➤ `pragma omp ... firstprivate(x)`

□  $x$  is a fundamental data type

- Private  $x$  is directly copied from the shared  $x$

□  $x$  is an array

- Copy the data with `sizeof(x)` to the private memory

□  $x$  is a pointer

- Private  $x$  points to the same location as the shared  $x$

□  $x$  is a class instance

- Copy constructor is called to create the private  $x$



## Clause: *firstprivate*

➤ `pragma omp ... firstprivate(x)`

□  $x$  is a fundamental data type

- Private  $x$  is directly copied from the shared  $x$

□  $x$  is an array

- Copy the data with `sizeof(x)` to the private memory

□  $x$  is a pointer

- Private  $x$  points to the same location as the shared  $x$

□  $x$  is a class instance

- Copy constructor is called to create the private  $x$

## Clause: *lastprivate*

- The *lastprivate* clause tells the compiler that the value of the private variable after the *sequentially last* loop iteration should be assigned to the shared variable upon loop exit
  - In other words, when the thread responsible for the sequentially last loop iteration exits the loop, its copy of the private variable is copied back to the shared variable





## *Example: lastprivate*

```
#pragma omp parallel for lastprivate (x)
```

```
for (i = 0; i < N; i++)
```


```
  x = foo(i);
```

```
  y[i] = bar(i, x);
```


```
}
```

```
last_x = x; // == foo(N-1)
```


# Work-Sharing Constructs



a type of "data  
parallelism".



a type of "functional  
parallelism".



serializes a section of  
code



# *Pragma: parallel*

- In the effort to increase grain size, sometimes the code that should be executed in parallel goes beyond a single *for* loop
  - ❑ The *parallel* pragma is used when a block of code should be executed in parallel
  - ❑ SPMD-style programming
  - ❑ Single program, multiple data

# *Pragma: for*

- The *for* pragma is used inside a block of code already marked with the *parallel* pragma
  - ❑ It indicates a *for* loop whose iterations should be divided among the active threads
  - ❑ There is a *barrier synchronization* of the threads at the end of the *for* loop

# *Pragma parallel and Pragma for*

## ◆ #pragma omp for

- Used inside a block of code already marked by *parallel*
- Distribute the iterations to the active threads

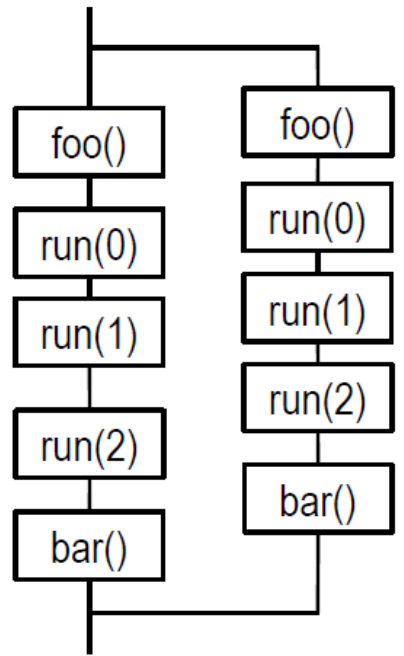
```
#pragma omp parallel \  
num_threads(2)  
{  
    foo();  
    #pragma omp for  
    for (int i = 0; i < 3; i++)  
        run(i);  
    bar();  
}
```

```
#pragma omp parallel \  
num_threads(2)  
{  
    foo();  
    for (int i = 0; i < 3; i++)  
        run(i);  
    bar();  
}
```

# Pragma parallel and Pragma for

## ◆ #pragma omp for

- Used inside a block of code already marked by *parallel*
- Distribute the iterations to the active threads



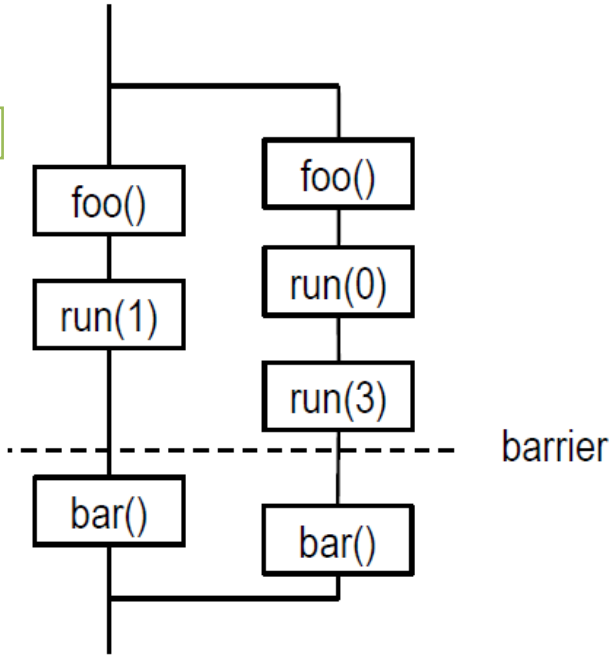
```
#pragma omp parallel \  
num_threads(2)  
{  
    foo();  
    for (int i = 0; i < 3; i++)  
        run(i);  
    bar();  
}
```

# Pragma parallel and Pragma for

- ◆ #pragma omp for
  - Used inside a block of code already marked by *parallel*
  - Distribute the iterations to the active threads

```
#pragma omp parallel \  
num_threads(2)  
{  
    foo();  
    #pragma omp for  
    for (int i = 0; i < 3; i++)  
        run(i);  
    bar();  
}
```

只能在并行结构中使用





# Pragma: *single*

- The *single* pragma is used inside a parallel block of code
  - ❑ It tells the compiler that **only a single thread** should execute the statement or block of code immediately following
  - ❑ May be useful when **dealing with sections of code that are not thread safe** (such as I/O)
  - ❑ Threads in the team that do **not execute the single directive**, **wait** at the end of the enclosed code block, **unless a nowait clause** is specified.

## *Example: master and single nowait*

```
tid = omp_get_thread_num();  
if (tid == 0) {  
    nthreads = omp_get_num_threads();  
    printf("Number of threads = %d\n", nthreads);  
}
```

=

```
#pragma omp master  
{  
    nthreads = omp_get_num_threads();  
    printf("Number of threads = %d\n", nthreads);  
}
```

≈

```
#pragma omp single nowait  
{  
    nthreads = omp_get_num_threads();  
    printf("Number of threads = %d\n", nthreads);  
}
```

## ***Pragma: sections and section***

- **Directive sections specifies that the enclosed section(s) of code are to be divided among the threads in the team.**
- **Independent section directives are nested within a sections directive.**
  - ❑ **Each section is executed once by a thread in the team.**
  - ❑ **Different sections may be executed by different threads.**
  - ❑ **It is possible for a thread to execute more than one section if it is quick enough and the implementation permits such.**

## *Example: sections and section*

```
#pragma omp parallel shared(a,b,c,d) private(i)
{
  #pragma omp sections
  {
    #pragma omp section
    {
      for (i=0; i<N; i++)
        c[i] = a[i] + b[i];
    }
    #pragma omp section
    {
      for (i=0; i<N; i++)
        d[i] = a[i] * b[i];
    }
  } /* end of sections */
} /* end of parallel section */
```

## Clause: reduction (归并)

- ◆ Reductions are so common that OpenMP provides a reduction clause for the *parallel*, *for*, and *sections*

#pragma omp ... reduction (op : list)

- A **private** copy of each list variable is created and initialized depending on the *op*
  - The identity value *op* (e.g., 0 for addition)
- These copies are **updated locally** by threads
- At end of construct, local copies are combined through *op* into a single value and combine the value in the original **shared** variable

# C/C++ Reduction Operation

结合率

- ◆ Reduction with an **associative** binary operator  $\oplus$

$$a_1 \oplus a_2 \oplus a_3 \oplus \dots \oplus a_n$$

- ◆ A range of associative and commutative operators can be used with reduction
- ◆ Initial values are the ones that make sense

Operator	Initial Value
+	0
*	1
-	0
^	0

Operator	Initial Value
&	~0
	0
&&	1
	0



# Reduction: an Artificial Example

```
int sum = 3;
int prod = 5;
#pragma omp parallel for \
    reduction(+:sum) \
    reduction(*:prod) \
    num_threads(2)
for (int i=0; i < 3; ++i) {
    int tid =
        omp_get_thread_num();
    sum |= i;
    prod += i;
    printf("thread(%d) "
        "sum=%d prod=%d\n",
        tid, sum, prod);
}
printf("results: "
    "sum=%d prod=%d\n",
    tid, sum, prod);
```

## ◆ Assume

- thread 0 executes the 1st and 2nd iterations, and
- thread 1 executes the 3<sup>rd</sup> iteration

## ◆ Possible outputs

```
thread(0)  sum=0  prod=1
thread(1)  sum=2  prod=3
thread(0)  sum=1  prod=2
results:  sum=6  prod=30
```

## Reduction: an Artificial Example

```
int sum = 3;
int prod = 5;
#pragma omp parallel for \
    reduction(+:sum) \
    reduction(*:prod) \
    num_threads(2)
for (int i=0; i < 3; ++i) {
    int tid =
        omp_get_thread_num();
    sum |= i;
    prod += i;
    printf("thread(%d) "
        "sum=%d prod=%d\n",
        tid, sum, prod);
}
printf("results: "
    "sum=%d prod=%d\n",
    tid, sum, prod);
```

◆ **initial**  $\text{sum}_{\text{private}} = 0$

▪ **for the reduction of +**

◆ **initial**  $\text{prod}_{\text{private}} = 1$

▪ **for the reduction of \***

◆ **At the end of “parallel  
for” with reduction**

▪  $\text{sum}_{\text{shared}} += \sum_{\text{thread}} \text{sum}_{\text{private}}$

▪  $\text{prod}_{\text{shared}} *= \prod_{\text{thread}} \text{prod}_{\text{private}}$



# ***Strengths and Weaknesses of OpenMP***

## **➤ Strengths**

- ❑ Incremental parallelization & sequential equivalence**
- ❑ Well-suited for domain decompositions**
- ❑ Available on \*nix and Windows**

## **➤ Weaknesses**

- ❑ Not well-tailored for functional decompositions**
- ❑ Compilers do not have to check for such errors as deadlocks  
and race conditions**



# Thank You !