



并行与分布式计算

Parallel & Distributed Computing

陈鹏飞
数据科学与计算机学院
2018-05-04



Lecture 8 — Programming with MPI

Pengfei Chen

School of Data and Computer Science

May 04, 2018




Outline:

- **Principles of Message-Passing Programming**
- **The Building Blocks: Send and Receive Operations**
- **MPI: the Message Passing Interface**
- **Overlapping Communication with Computation**
- **Groups and Communicators**

A. Grama et al., “Introduction to Parallel Computing,” Addison Wesley, 2003

Message-Passing Model



Principles of Message-Passing Programming

- The logical view of a machine supporting the message passing paradigm consists of p processes, each with its own exclusive address space
- **CONSTRAINTS (限制) ...**
 - ❑ Each data element must belong to one of the partitions of the space; hence, data must be explicitly partitioned and placed
 - ❑ All interactions (read-only or read/write) require cooperation of two processes - the process that has the data and the process that wants to access the data
- These two constraints, while onerous (繁重), make underlying costs very explicit to the programmer

Principles of Message-Passing Programming

- Message-passing programs are often written using the *asynchronous* or *loosely synchronous* paradigms
 - ❑ In the asynchronous paradigm, all concurrent tasks execute asynchronously
 - ❑ In the loosely synchronous model, tasks or subsets of tasks synchronize to perform interactions. Between these interactions, tasks execute completely asynchronously
- Most message-passing programs are written using the *single program multiple data (SPMD)* model

The Building Blocks: Send and Receive Operations

- The prototypes of these operations are as follows:

```
send(void *sendbuf, int nelems, int dest)
```

```
receive(void *recvbuf, int nelems, int source)
```

- Consider the following code segments:

P0

```
a = 100;
```

```
send(&a, 1, 1);
```

```
a = 0;
```

P1

```
receive(&a, 1, 0)
```

```
printf("%d\n", a);
```

- ❑ The semantics of the send operation require that the value received by process **P1** must be *100* as opposed to *0*

- This motivates the design of the send and receive protocols



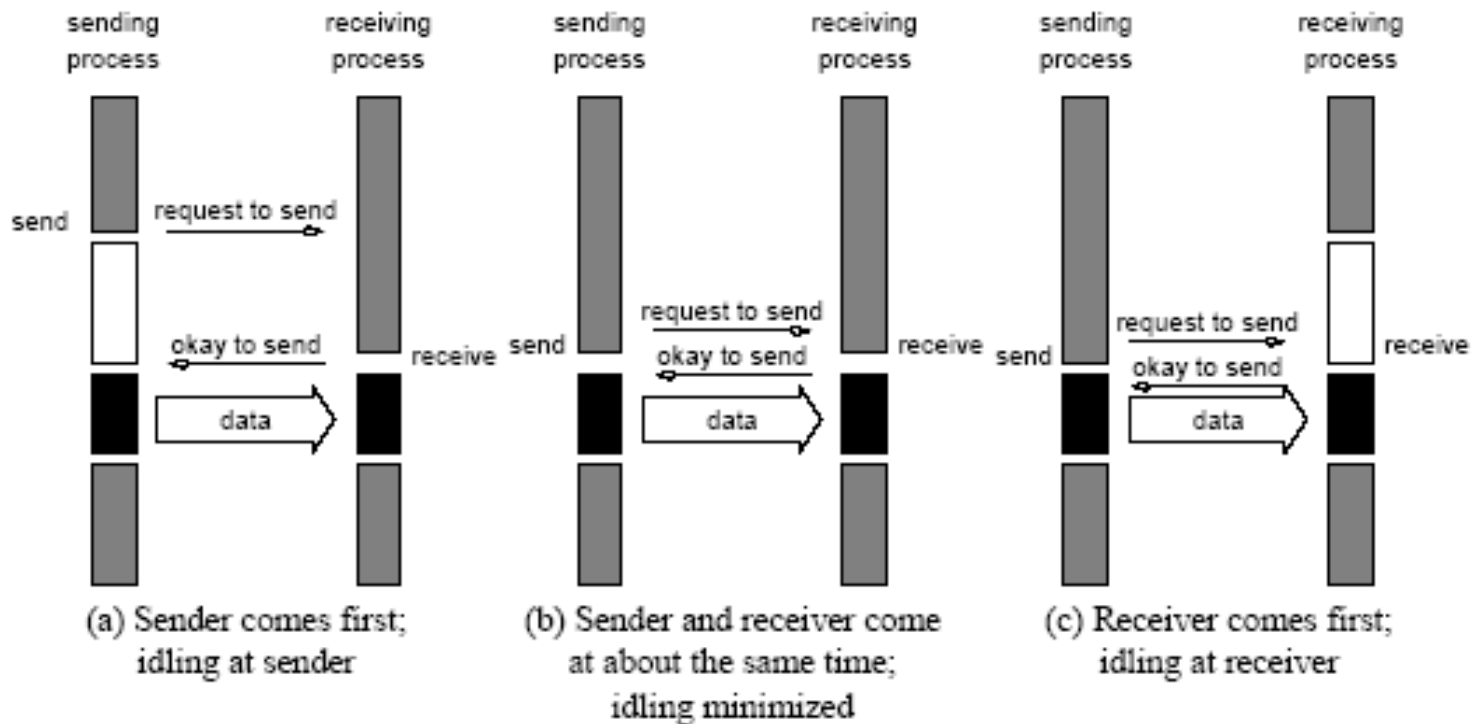
Send and Receive Operations

- **Non-buffered blocking**
- **Buffered blocking**
- **Non-blocking**

Non-Buffered Blocking Message Passing Operations

- A simple method for forcing (强制) send/receive semantics is for the send operation to return only when it is safe to do so
- In the non-buffered blocking send, the operation does not return until the matching receive has been encountered at the receiving process
- Idling and deadlocks are major issues with nonbuffered blocking sends

Non-Buffered Blocking Message Passing Operations



Handshake for a blocking non-buffered send/receive operation.
It is easy to see that in cases where sender and receiver do not reach communication point at similar times, there can be considerable idling overheads

Non-Buffered Blocking \Rightarrow Buffered Blocking

- In buffered blocking sends, the sender simply copies the data into the designated buffer and returns after the copy operation has been completed. The data is copied at a buffer at the receiving end as well
- Buffering alleviates idling at the expense of copying overheads




Non-Buffered Blocking \Rightarrow Buffered Blocking

- In buffered blocking sends, the sender simply copies the data into the designated buffer and returns after the copy operation has been completed. The data is copied at a buffer at the receiving end as well
- Buffering alleviates idling at the expense of copying overheads

Buffered Blocking Message Passing Operations

- In buffered blocking sends, the sender simply copies the data into the designated buffer and returns after the copy operation has been completed.
- The data is copied at a buffer at the receiving end as well
- Buffering trades off idling overhead for buffer copying overhead

Buffered Blocking Message Passing Operations



Blocking buffered transfer protocols: (a) in the *presence* of communication hardware with buffers at send and receive ends; and (b) in the *absence* of communication hardware, sender interrupts receiver and deposits data in buffer at receiver end.

Buffered Blocking Message Passing Operations

- **Bounded buffer sizes can have significant impact on performance**

P0

```
for (i = 0; i < 1000; i++){  
    produce_data(&a);  
    send(&a, 1, 1);  
}
```

P1

```
for (i = 0; i < 1000; i++){  
    receive(&a, 1, 0);  
    consume_data(&a);  
}
```

What if consumer was much slower than producer?

A. Grama et al., “Introduction to Parallel Computing,” Addison Wesley, 2003

Buffered Blocking Message Passing Operations

- **Bounded buffer sizes can have significant impact on performance**

P0

```
for (i = 0; i < 1000; i++){  
    produce_data(&a);  
    send(&a, 1, 1);  
}
```

P1

```
for (i = 0; i < 1000; i++){  
    receive(&a, 1, 0);  
    consume_data(&a);  
}
```

What if consumer was much slower than producer?

A. Grama et al., “Introduction to Parallel Computing,” Addison Wesley, 2003

Buffered Blocking Message Passing Operations

- **Bounded buffer sizes can have significant impact on performance**

P0

```
for (i = 0; i < 1000; i++){  
    produce_data(&a);  
    send(&a, 1, 1);  
}
```

P1

```
for (i = 0; i < 1000; i++){  
    receive(&a, 1, 0);  
    consume_data(&a);  
}
```

What if consumer was much slower than producer?

A. Grama et al., “Introduction to Parallel Computing,” Addison Wesley, 2003

Buffered Blocking Message Passing Operations

- Deadlocks are still possible with buffering since receive operations block.

P0

```
receive(&b, 1, 1);
```

```
send(&a, 1, 1);
```

P1

```
receive(&a, 1, 0);
```

```
send(&b, 1, 0);
```

Buffered Blocking Message Passing Operations

- Deadlocks are still possible with buffering since receive operations block.

P0

```
receive(&b, 1, 1);
```

```
send(&a, 1, 1);
```

P1

```
receive(&a, 1, 0);
```

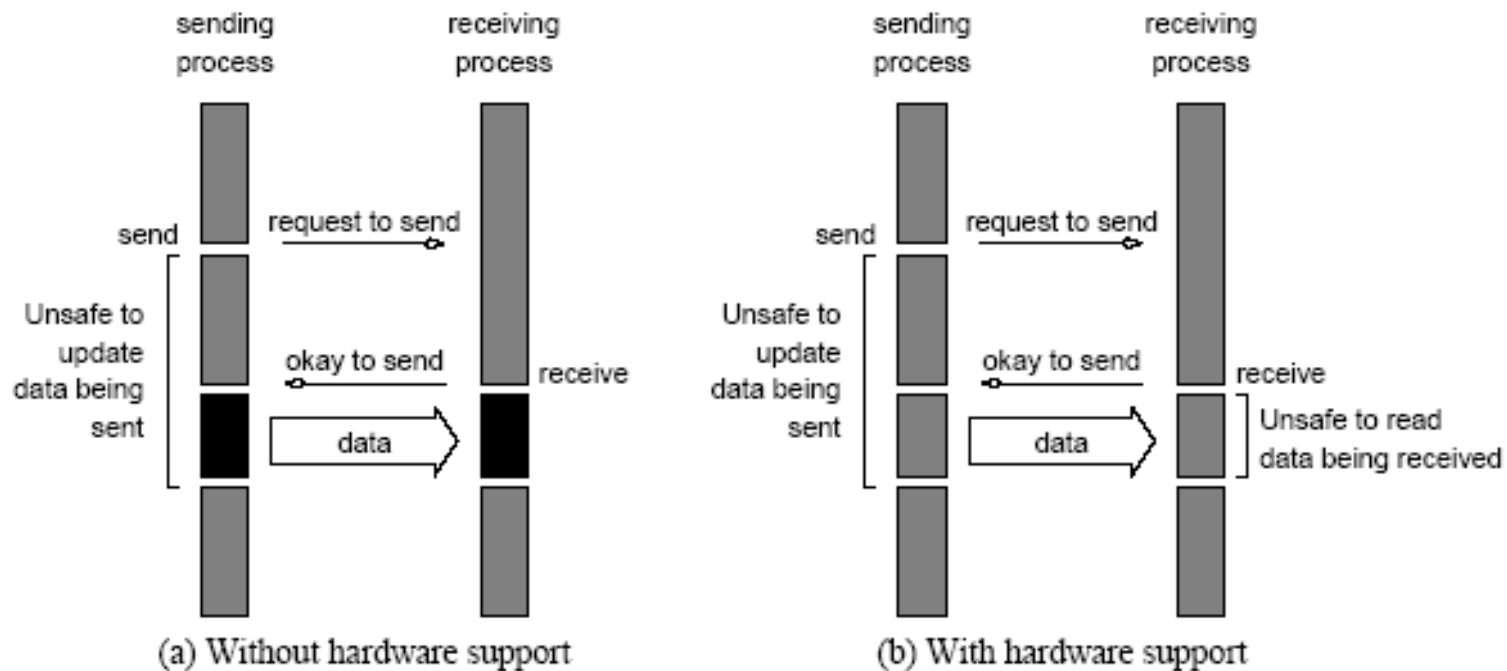
```
send(&b, 1, 0);
```

Non-Blocking Message Passing Operations

- **The programmer must ensure semantics of the send and receive.**
- **This class of non-blocking protocols returns from the send or receive operation before it is semantically safe to do so.**
- **Non-blocking operations are generally accompanied by a check-status operation.**
- **When used correctly, these primitives are capable of overlapping communication overheads with useful computations.**
- **Message passing libraries typically provide both blocking and non-blocking primitives.**

A. Grama et al., “Introduction to Parallel Computing,” Addison Wesley, 2003

Non-Blocking Message Passing Operations



Non-blocking non-buffered send and receive operations (a) in *absence* of communication hardware; (b) in *presence* of communication hardware.

Send and Receive Protocols

	Blocking Operations	Non-Blocking Operations
Buffered	<p>Sending process returns after data has been copied into communication buffer</p>	<p>Sending process returns after initiating DMA transfer to buffer. This operation may not be completed on return</p>
Non-Buffered	<p>Sending process blocks until matching receive operation has been encountered</p> <p>Send and Receive semantics assured by corresponding operation</p>	<p>Programmer must explicitly ensure semantics by polling to verify completion</p>

Space of possible protocols for send and receive operations.

A. Grama et al., “Introduction to Parallel Computing,” Addison Wesley, 2003

The Message Passing Interface

- **Late 1980s: vendors had unique libraries**
- **1989: Parallel Virtual Machine (PVM) developed at Oak Ridge National Lab**
- **1992: Work on MPI standard begun**
- **1994: Version 1.0 of MPI standard**
- **1997: Version 2.0 of MPI standard**
- **Today: MPI is dominant message passing library standard**

MPI: the Message Passing Interface

- MPI defines a standard *library* for message-passing that can be used to develop portable message-passing programs using either C or Fortran.
- The MPI standard defines both the syntax as well as the semantics of a core set of library routines.
- Vendor implementations of MPI are available on almost all commercial parallel computers.
- It is possible to write fully-functional message-passing programs by using only the six routines.

A. Grama et al., “Introduction to Parallel Computing,” Addison Wesley, 2003



MPI Implementations and Tutorials

➤ **Standard**

❑ <http://www.mpi-forum.org>

➤ **Implementations**

❑ MPICH2 <http://www.mcs.anl.gov/research/projects/mpich2/>

❑ Open MPI <http://www.open-mpi.org/>

➤ **Tutorials**

❑ <https://computing.llnl.gov/tutorials/mpi/>

❑ <http://www.mcs.anl.gov/research/projects/mpi/>

MPI: the Message Passing Interface

The minimal set of MPI routines.

<code>MPI_Init</code>	Initializes MPI.
<code>MPI_Finalize</code>	Terminates MPI.
<code>MPI_Comm_size</code>	Determines the number of processes.
<code>MPI_Comm_rank</code>	Determines the label of calling process.
<code>MPI_Send</code>	Sends a message.
<code>MPI_Recv</code>	Receives a message.

Starting and Terminating the MPI Library

- *MPI_Init* is called prior to any calls to other MPI routines. Its purpose is to **initialize the MPI environment**
- *MPI_Finalize* is called at the end of the computation, and it performs various **clean-up tasks to terminate the MPI environment**
- **The prototypes of these two functions are:**


```
int MPI_Init(int *argc, char ***argv)
int MPI_Finalize()
```
- *MPI_Init* **also strips off any MPI related command-line arguments.**
- **All MPI routines, data-types, and constants are prefixed by “*MPI_*”. The return code for successful completion is *MPI_SUCCESS***

Some *Basic Concepts*

- Processes can be collected into groups
- Each message is sent in a context, and must be received in the same context
 - Provides necessary support for libraries
- A group and context together form a communicator
- A process is identified by its rank in the group associated with a communicator
- There is a default communicator whose group contains all initial processes, called *MPI_COMM_WORLD*

A. Grama et al., “Introduction to Parallel Computing,” Addison Wesley, 2003

Groups and Communicators



➤ Related MPI functions

- ❑ Form new group as a subset of global group using *`MPI_Group_incl`*
- ❑ Create new communicator for new group using *`MPI_Comm_create`*
- ❑ Determine new rank in new communicator using *`MPI_Comm_rank`*
- ❑ Conduct communications using any MPI message passing routine
- ❑ When finished, free up new communicator and group (optional) using *`MPI_Comm_free`* and *`MPI_Group_free`*

Communicators

- A communicator defines a *communication domain* - a set of processes that are allowed to communicate with each other
- Information about communication domains is stored in variables of type `MPI_Comm`
- Communicators are used as arguments to all message transfer MPI routines
- A process can belong to many different (possibly overlapping) communication domains
- MPI defines a default communicator called `MPI_COMM_WORLD` which includes all the processes

Querying Information

- The `MPI_Comm_size` and `MPI_Comm_rank` functions are used to determine the number of processes and the label of the calling process, respectively.

- The calling sequences of these routines are as follows:

```
int MPI_Comm_size(MPI_Comm comm, int *size)
```

```
int MPI_Comm_rank(MPI_Comm comm, int *rank)
```

- The rank of a process is an integer that ranges from zero up to the size of the communicator minus one

MPI: Hello World!

```
#include <mpi.h>

main(int argc, char *argv[])
{
    int npes, myrank;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &npes);
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
    printf("From process %d out of %d, Hello World!\n",
           myrank, npes);
    MPI_Finalize();
}
```



Compiling and Running MPI Programs

➤ Compiling examples

- ❑ `mpicc -o foo foo.c`

- ❑ `mpic++ -o bar bar.cpp`

➤ Running examples

- ❑ `mpirun -np 4 foo`

- ❑ `mpirun -np 2 foo : -np 4 bar`

➤ Specifying host processors

- ❑ see “`--hostfile`” and “`--host`” options



MPI Messages

- **data : (address, count, datatype)**
 - ❑ **Datatype is either a predefined type, or a custom type**
- **message : (data, tag)**
 - ❑ **Tag is an integer to assist the receiving process in identifying the message**

Sending and Receiving Messages

- **The basic functions for sending and receiving messages in MPI are the *MPI_Send* and *MPI_Recv*, respectively**
- **The calling sequences of these routines are as follows:**

```
int MPI_Send(void *buf, int count, MPI_Datatype
datatype, int dest, int tag, MPI_Comm comm)
int MPI_Recv(void *buf, int count, MPI_Datatype
datatype, int source, int tag,
MPI_Comm comm, MPI_Status *status)
```
- **MPI provides equivalent datatypes for all C datatypes. This is done for portability reasons**
- **The datatype *MPI_BYTE* corresponds to a byte (8 bits) and *MPI_PACKED* corresponds to a collection of data items that has been created by packing non-contiguous data**
- **The message-tag can take values ranging from zero up to the MPI defined constant *MPI_TAG_UB***

MPI's Send Modes (1/2)

- **MPI_Send**
 - Will not return until you can use the send buffer
- **MPI_Bsend**
 - Returns immediately and you can use the send buffer
 - Related: `MPI_buffer_attach()`, `MPI_buffer_detach()`
- **MPI_Ssend**
 - Will not return until matching receive posted
 - Send + synchronous communication semantics
- **MPI_Rsend**
 - May be used **ONLY** if matching receive already posted
 - The sender provides additional information to the system that can save some overhead

MPI's Send Modes (2/2)

- **MPI_Isend**
 - Nonblocking send, but you can **NOT** reuse the send buffer immediately
 - Related: **MPI_Wait()**, **MPI_Test()**
- **MPI_Ibsend**
- **MPI_Issend**
- **MPI_Irsend**

MPI Datatypes

MPI Datatype	C Datatype
MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	
MPI_PACKED	

A. Grama et al., “Introduction to Parallel Computing,” Addison Wesley, 2003



Sending and Receiving Messages

- **MPI allows specification of wildcard arguments for both source and tag.**
- **If source is set to `MPI_ANY_SOURCE`, then any process of the communication domain can be the source of the message.**
- **If tag is set to `MPI_ANY_TAG`, then messages with any tag are accepted.**
- **On the receive side, the message must be of length equal to or less than the length field specified.**

A. Grama et al., "Introduction to Parallel Computing," Addison Wesley, 2003

Avoiding Deadlocks

Consider:

```
int a[10], b[10], myrank;
MPI_Status status;
...
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
if (myrank == 0) {
    MPI_Send(a, 10, MPI_INT, 1, 1, MPI_COMM_WORLD);
    MPI_Send(b, 10, MPI_INT, 1, 2, MPI_COMM_WORLD);
}
else if (myrank == 1) {
    MPI_Recv(b, 10, MPI_INT, 0, 2, MPI_COMM_WORLD);
    MPI_Recv(a, 10, MPI_INT, 0, 1, MPI_COMM_WORLD);
}
...
```

If *MPI_Send* is blocking, there is a deadlock.

Sending and Receiving Messages

- **On the receiving end, the status variable can be used to get information about the *MPI_Recv* operation**

- **The corresponding data structure contains:**

```
typedef struct MPI_Status {  
  
    int MPI_SOURCE;  
  
    int MPI_TAG;  
  
    int MPI_ERROR; };
```

- **The *MPI_Get_count* function returns the precise count of data items received**

```
int MPI_Get_count(MPI_Status *status, MPI_Datatype datatype, int *count)
```



Thank You !