



并行与分布式计算

Parallel & Distributed Computing

陈鹏飞
数据科学与计算机学院
2018-05-25
chenpf7@mail.sysu.edu.cn



Lecture 11 — Parallel Computing with MapReduce

Pengfei Chen

School of Data and Computer Science

May 25, 2018

chenpf7@mail.sysu.edu.cn



Outline:

- MapReduce Programming Model
- Typical Problems Solved by MapReduce
- MapReduce Examples
- A Brief History
- MapReduce Execution Overview
- Hadoop



Motivation: Large Scale Data Processing

- Want to process lots of data (>1TB)
- Want to parallelize across hundreds/thousands of CPUs
- ... Want to make this easy



Motivation: Large Scale Data Processing

- Want to process lots of data (>1TB)
- Want to parallelize across hundreds/thousands of CPUs
- ... Want to make this easy



MapReduce

- A simple and powerful interface that enables automatic parallelization and distribution of large-scale computations, combined with an implementation of this interface that achieves high performance on large clusters of commodity PCs.”
- **More simply**, MapReduce is A parallel programming model and associated implementation

Dean and Ghemawat, “MapReduce: Simplified Data Processing on Large Clusters”, Google Inc.

Some MapReduce Terminology

- Job—A “full program” -an execution of a Mapper and Reducer across a data set
- Task —An execution of a Mapper or a Reducer on a slice of data
 - ▣ a.k.a. Task-In-Progress (TIP)
- Task Attempt —A particular instance of an attempt to execute a task on a machine



Terminology Example

- Running “Word Count” across 20 files is one job
- 20 files to be mapped imply 20 map tasks+ some number of reduce tasks
- At least 20 map task attemptswill be performed... more if a machine crashes, etc.



Task Attempts

- A particular task will be attempted at least once, possibly more times if it crashes
 - ❑ If the same input causes crashes over and over, that input will eventually be abandoned
- Multiple attempts at one task may occur in parallel with speculative execution turned on
 - ❑ Task ID from TaskInProgress is not a unique identifier



MapReduce Programming Model

- Process data using special `map()` and `reduce()` functions
 - ❑ The `map()` function is called on every item in the input and emits a series of intermediate key/value pairs
 - ❑ All values associated with a given key are grouped together
 - ❑ The `reduce()` function is called on every unique key, and its value list, and emits a value that is added to the output



Map

- Records from the data source (lines out of files, rows of a database, etc) are fed into the map function as key*value pairs: e.g., (filename, line)
- map() produces one or more intermediate values along with an output key from the input

```
- map (in_key, in_value) ->  
      (out_key, intermediate_value) list
```


reduce

- After the map phase is over, all the intermediate values for a given output key are combined together into a list
- `reduce()` combines those intermediate values into one or more final values for that same output key

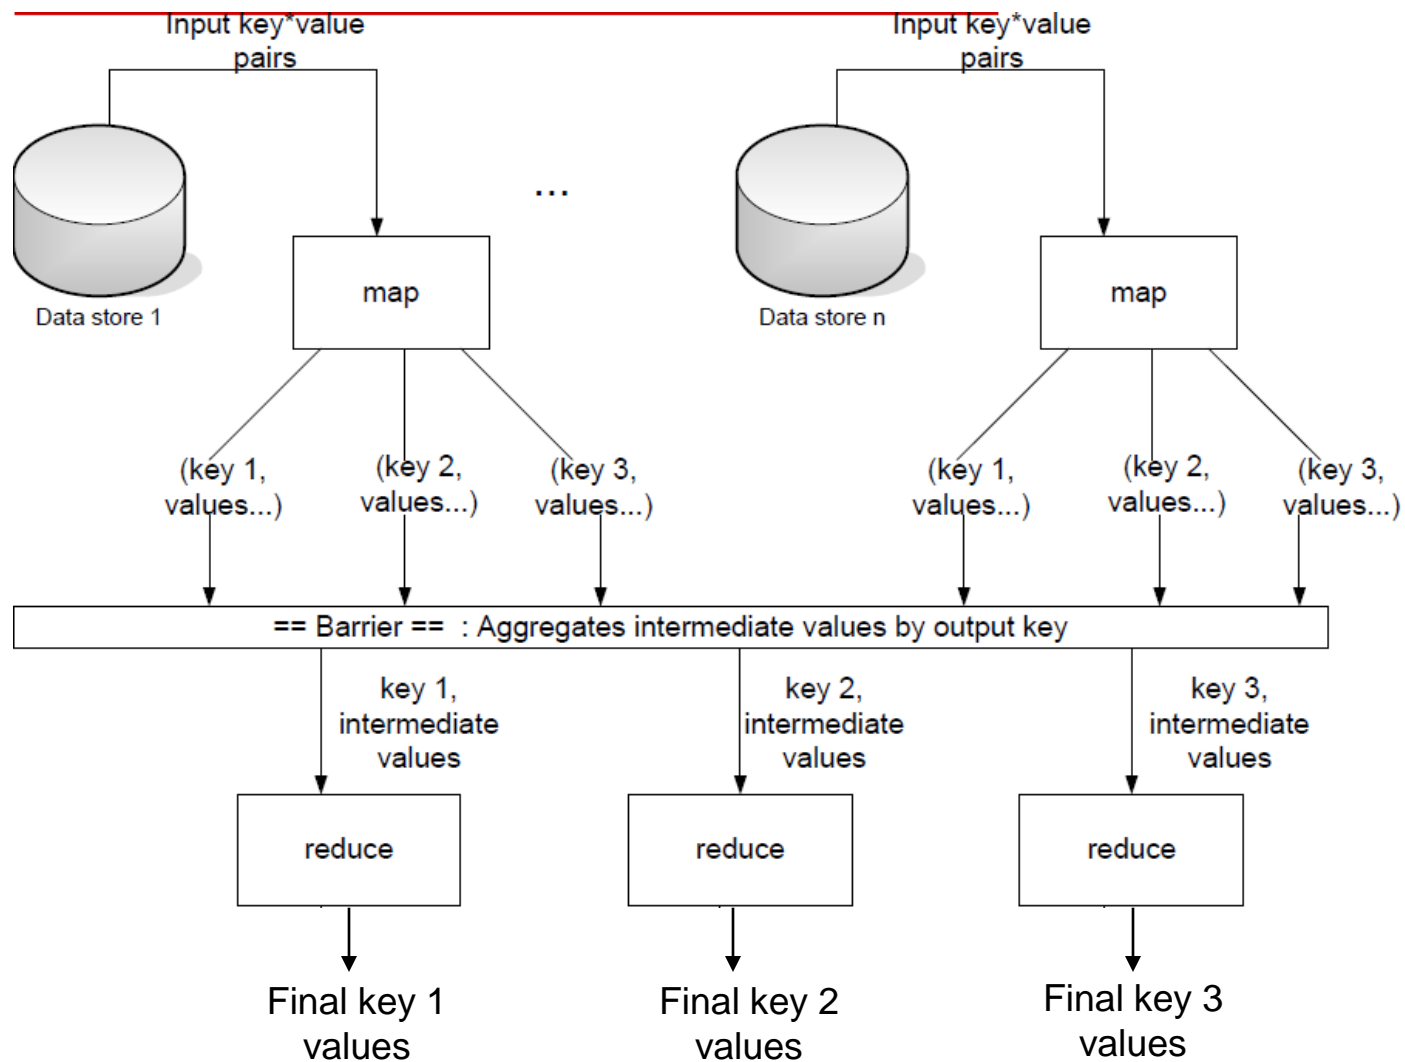
```
- reduce (out_key, intermediate_value list) ->  
  out_value list
```

reduce

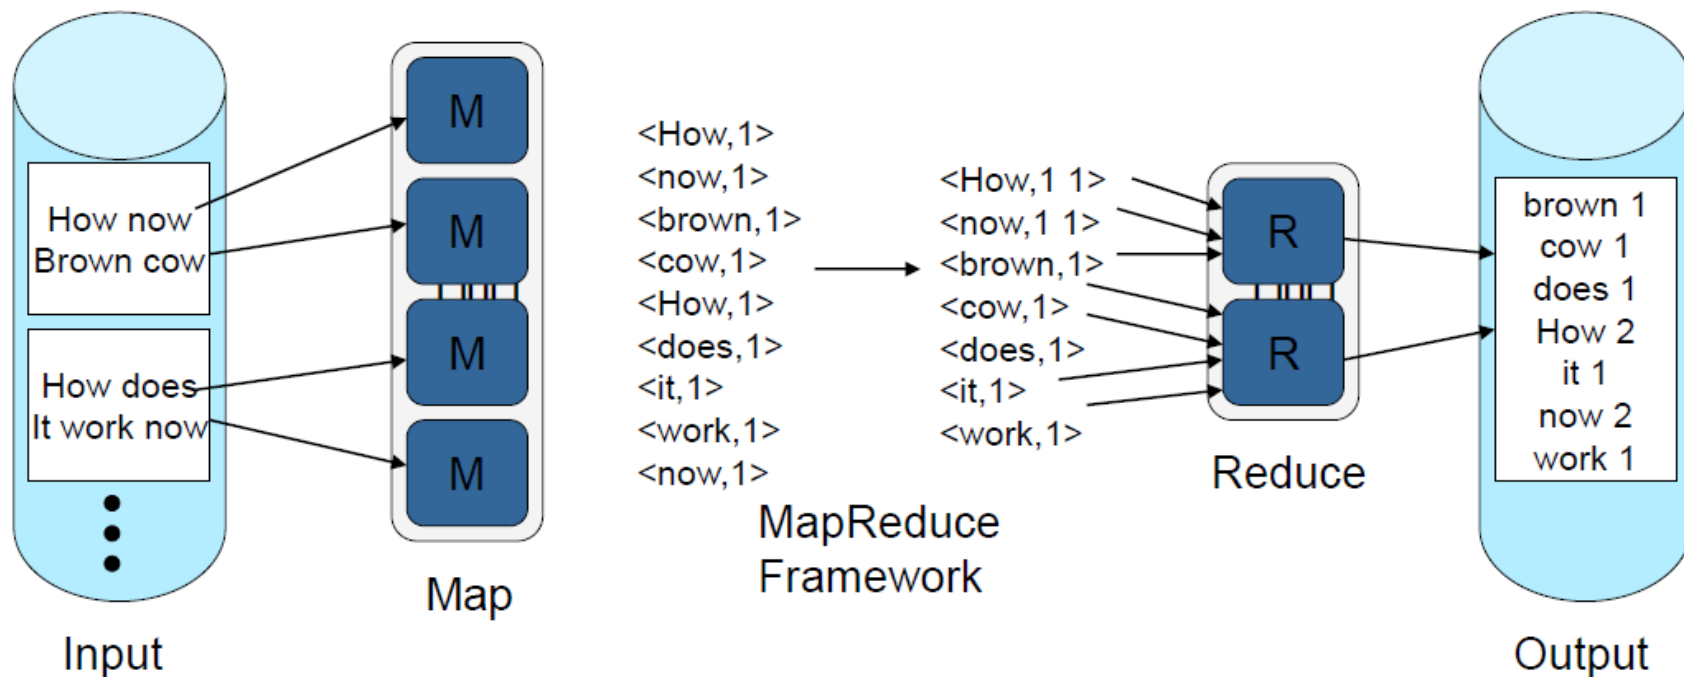
```
reduce (out_key, intermediate_value list) ->  
      out_value list
```



MapReduce Architecture



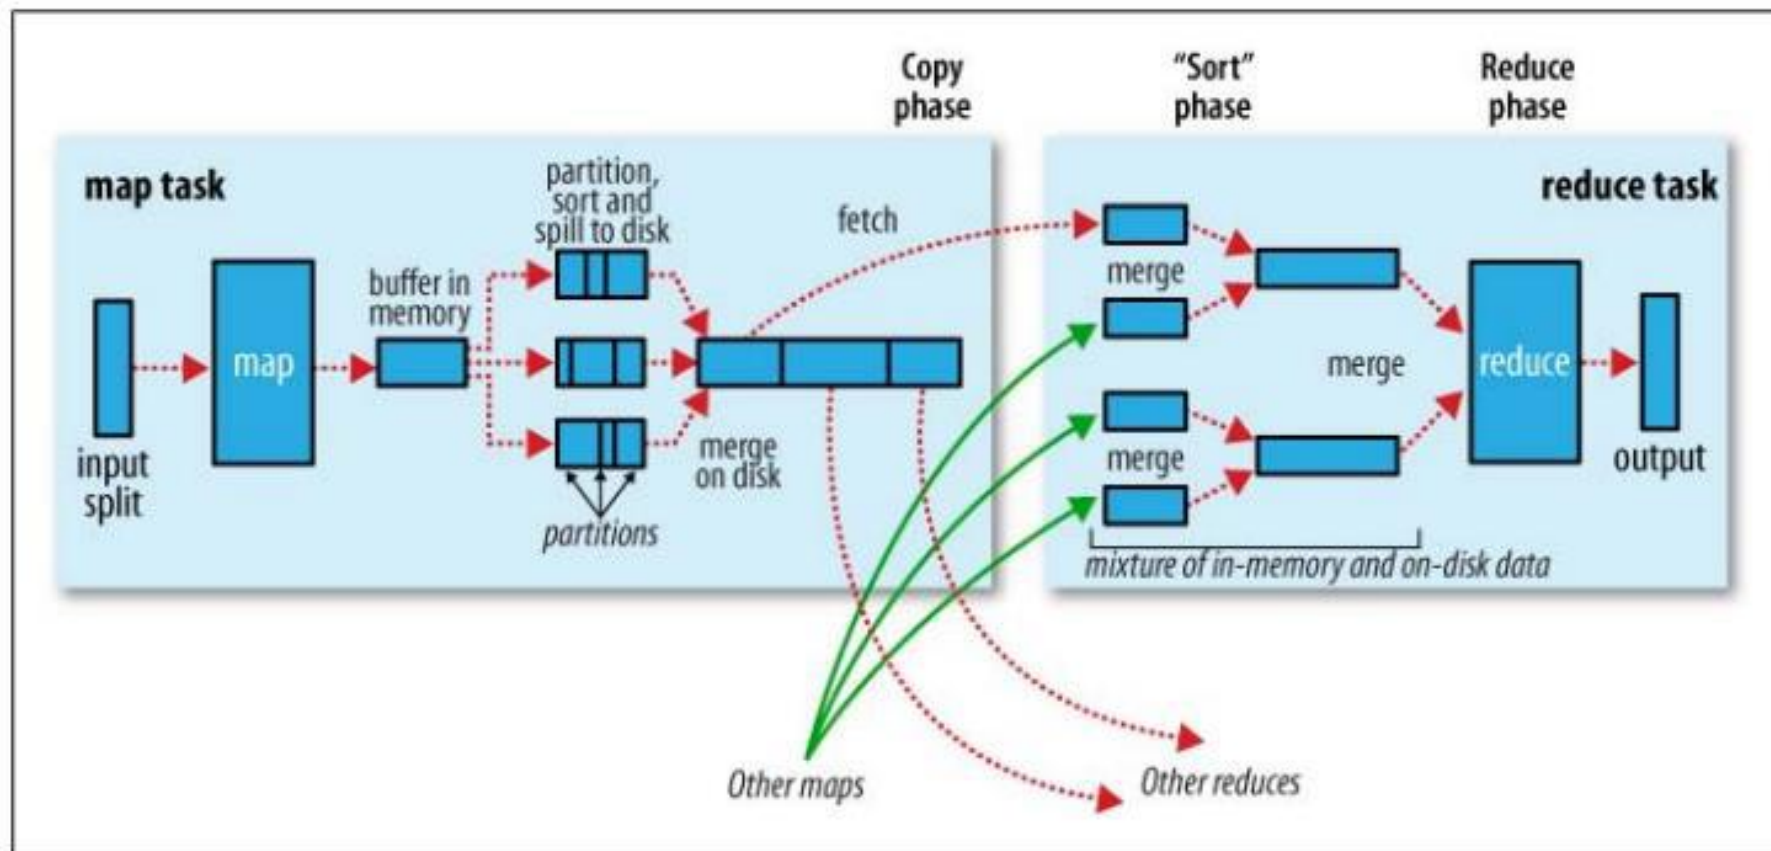
MapReduce Programming Model



□ More formally,

- $\text{Map}(k1, v1) \rightarrow \text{list}(k2, v2)$
- $\text{Reduce}(k2, \text{list}(v2)) \rightarrow \text{list}(v2)$

MapReduce in One Picture





MapReduce Runtime System

- Partitions input data
- Schedules execution across a set of machines
- Handles machine failure
- Manages interprocess communication



Parallelism

- `map()` functions run in parallel, creating different intermediate values from different input data sets
- `reduce()` functions also run in parallel, each working on a different output key
- All values are processed independently
- Bottleneck: reduce phase can't start until map phase is completely finished



Locality

- Master program divides up tasks based on location of data: tries to have map() tasks on same machine as physical file data, or at least same rack
- map() task inputs are divided into 64 MB blocks: same size as Google File System chunks




Fault Tolerance

- Master detects worker failures
 - ❑ Re-executes completed & in-progress map() tasks
 - ❑ Re-executes in-progress reduce() tasks
- Master notices particular input key/values cause crashes in map(), and skips those values on re-execution
 - ❑ Effect: Can work around bugs in third-party libraries!

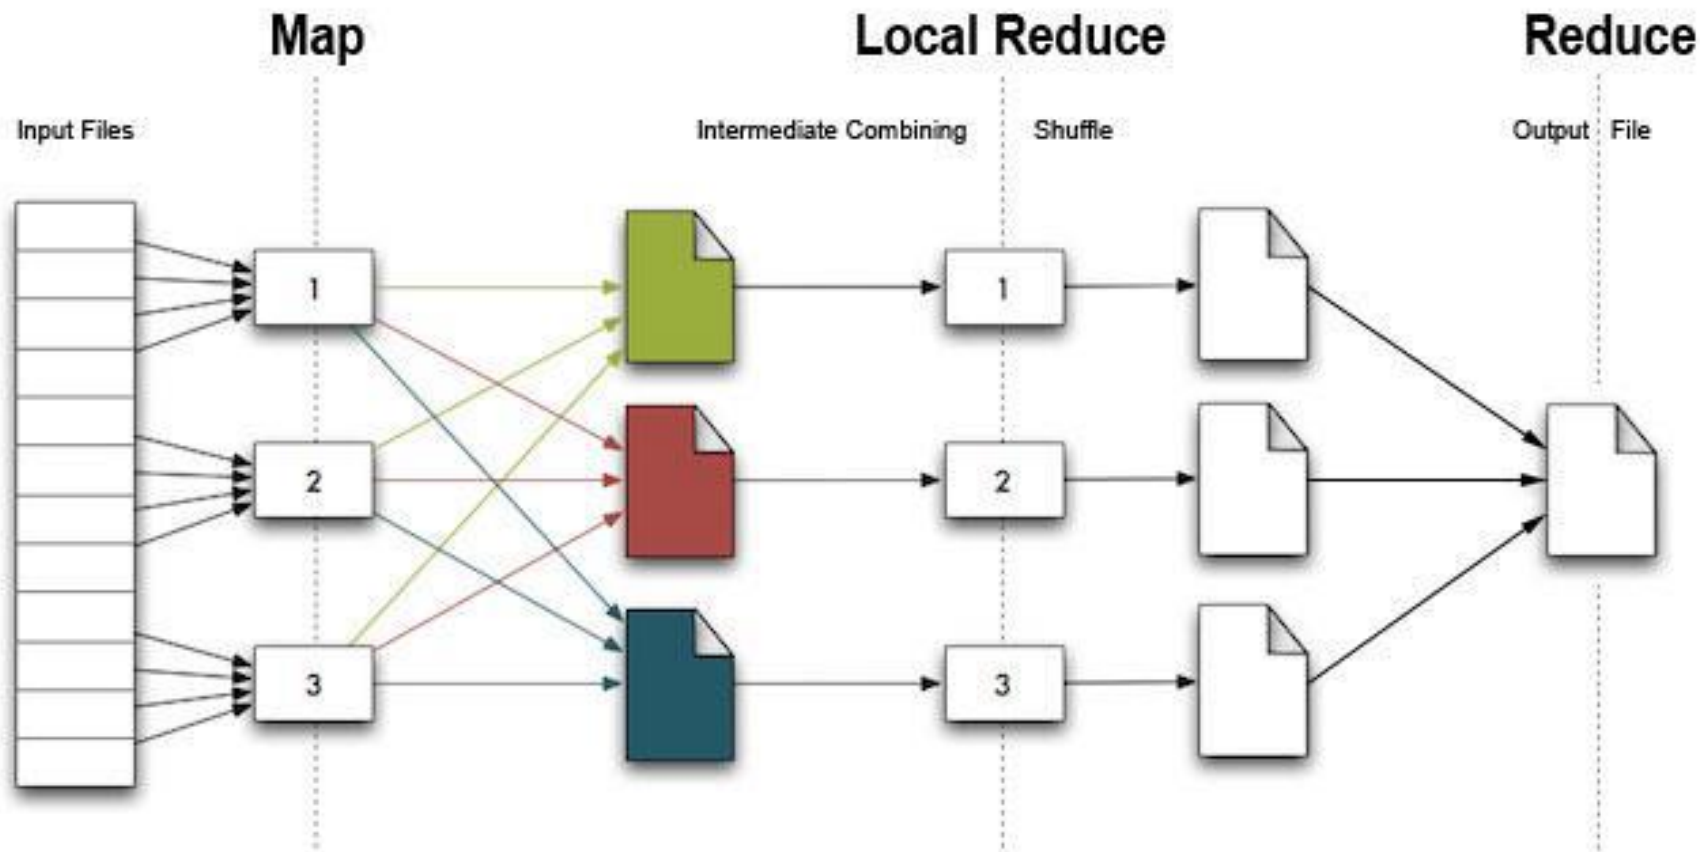
Optimizations

- No reduce can start until map is complete
 - ❑ A single slow disk controller can rate-limit the whole process
- Master redundantly executes “slow-moving” map tasks; uses results of first copy to finish
- “Combiner” functions can run on same machine as a mapper
- Causes a mini-reduce phase to occur before the real reduce phase, to save bandwidth

Combiner



Optimizations





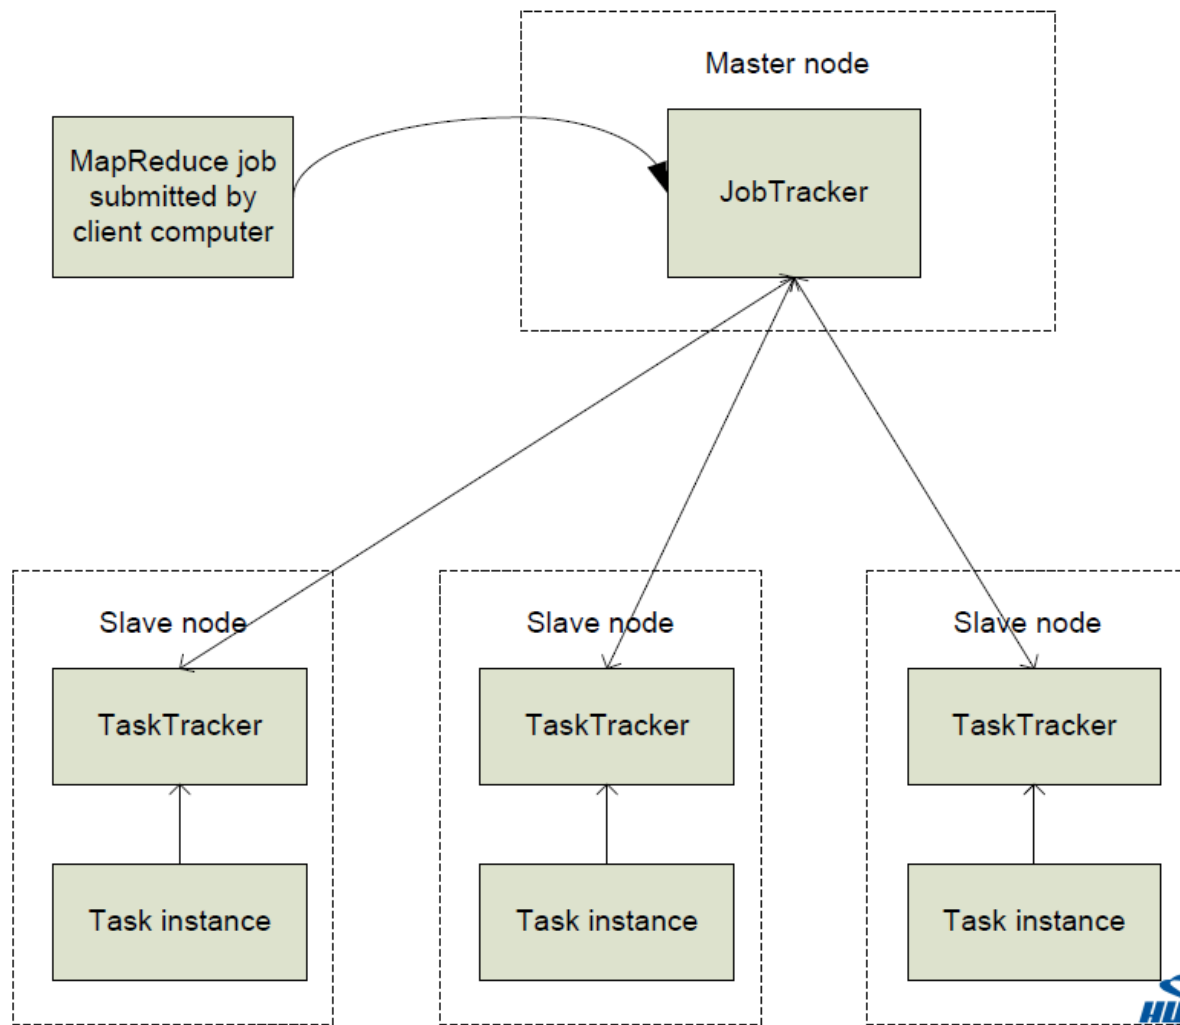
MapReduce Benefits

- Greatly reduces parallel programming complexity
 - Reduces synchronization complexity
 - Automatically partitions data
 - Provides failure transparency
 - Handles load balancing



Typical Problems Solved by MapReduce

MapReduce: High Level





Nodes, Trackers, Tasks

- Master node runs JobTrackerinstance, which accepts Job requests from clients
- TaskTracker instances run on slave nodes
- TaskTrackerforks separate Java process for task instances

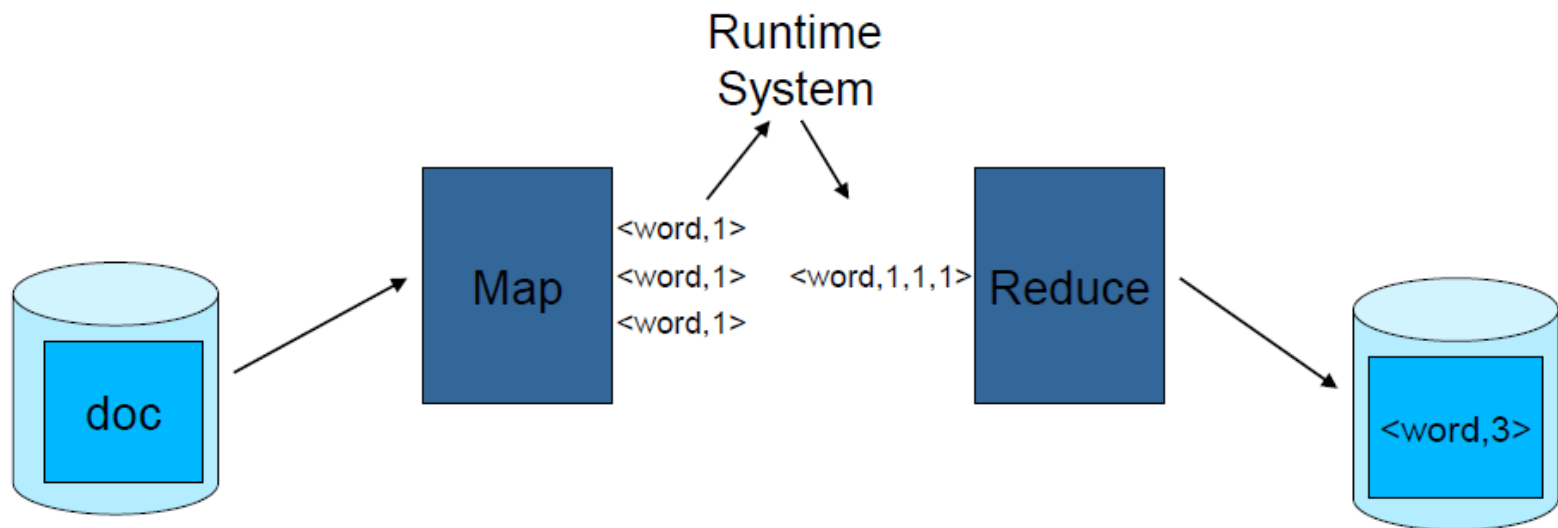
Typical Problems Solved by MapReduce

- ☐ Read a lot of data
 - ☐ **Map**: extract something you care about from each record
 - ☐ **Shuffle** and **Sort**
 - ☐ **Reduce**: aggregate, summarize, filter, or transform
 - ☐ Write the results
-
- ☐ Outline stays the same, but **map** and **reduce** change to fit the problem

MapReduce Examples

MapReduce Examples

➤ Word frequency



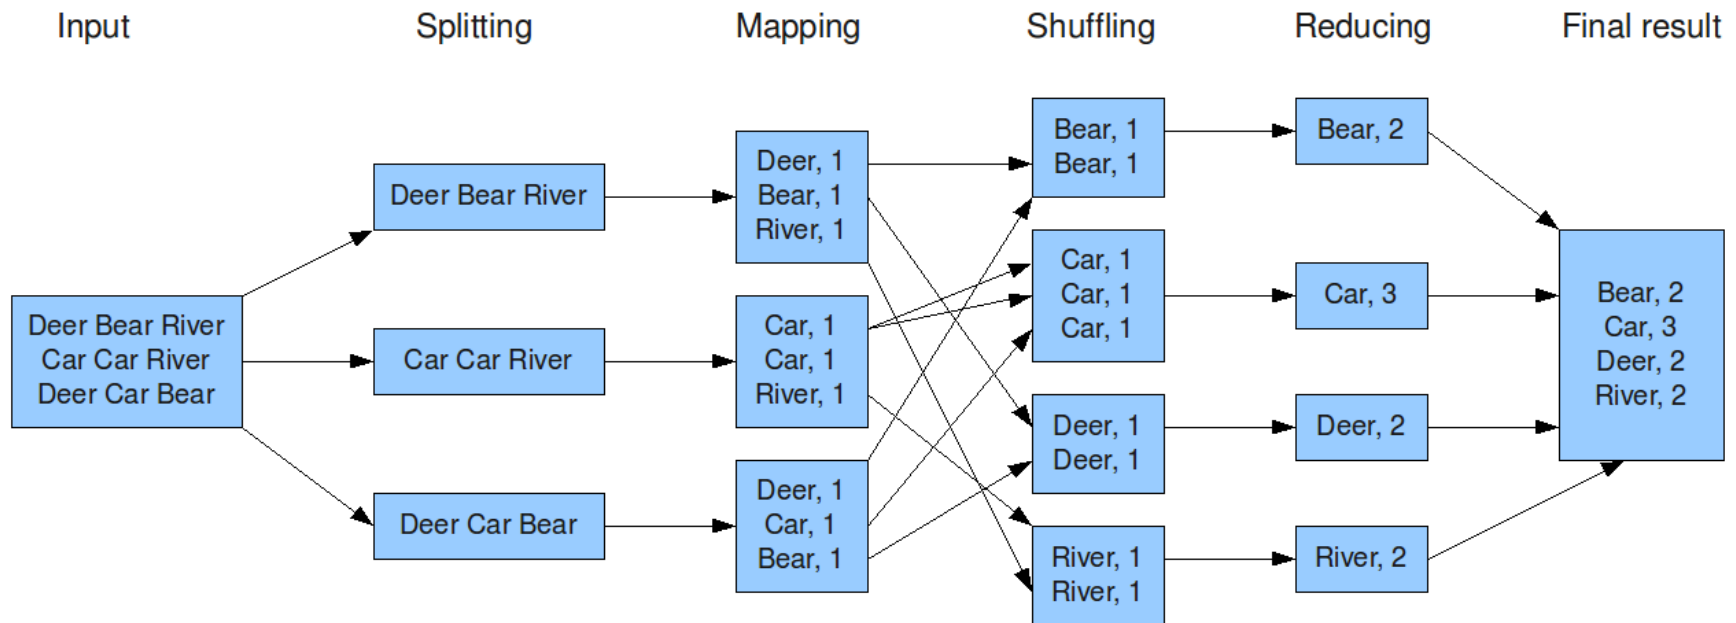


Example: Count Word Occurrences

```
map(String input_key, String input_value):  
    // input_key: document name  
    // input_value: document contents  
    for each word w in input_value:  
        EmitIntermediate(w, "1");  
  
reduce(String output_key, Iterator  
    intermediate_values):  
    // output_key: a word  
    // output_values: a list of counts  
    int result = 0;  
    for each v in intermediate_values:  
        result += ParseInt(v);  
    Emit(AsString(result));
```

Example: Count Word Occurrences

The overall MapReduce word count process



MapReduce Examples

□ Distributed grep

- Map function emits `<word, line_number>` if word matches search criteria
- Reduce function is the identity function

□ URL access frequency

- Map function processes web logs, emits `<url, 1>`
- Reduce function sums values and emits `<url, total>`



A Brief History



A Brief History

MapReduce is a new use of an old idea in Computer Science

➤ Map: Apply a function to every object in a list

□ Each object is independent

- Order is unimportant
- Maps can be done in parallel

□ The function produces a result

➤ Reduce: Combine the results to produce a final result

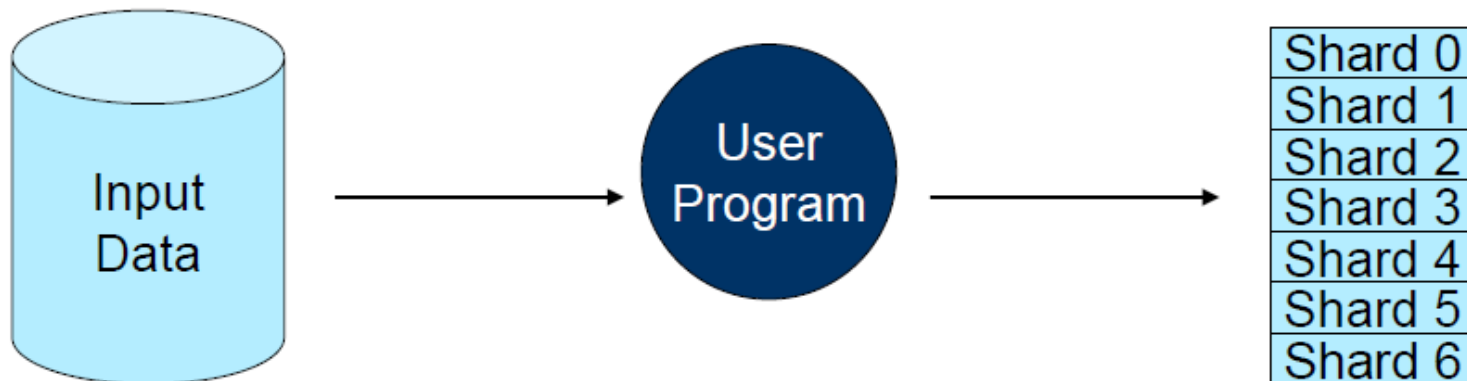
You may have seen this in a Lisp or functional programming course



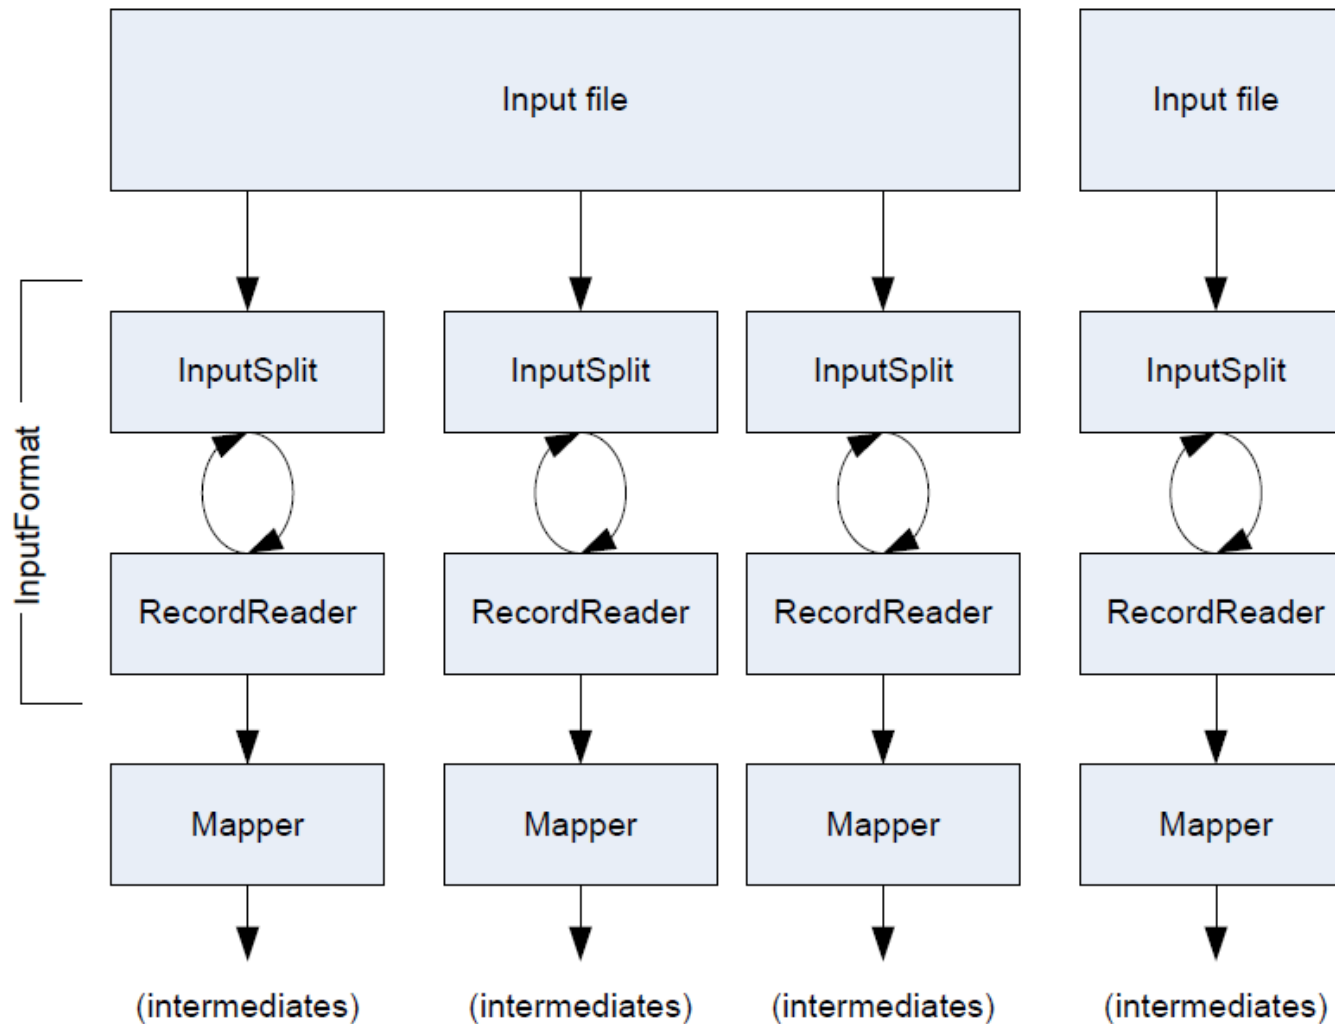
MapReduce Execution Overview

MapReduce Execution Overview

- The user program, via the MapReduce library, shards the input data




Getting Data To The Mapper



Getting Data To The Mapper


- The user program creates process copies distributed on a machine cluster. One copy will be the “master” and the others will be worker threads



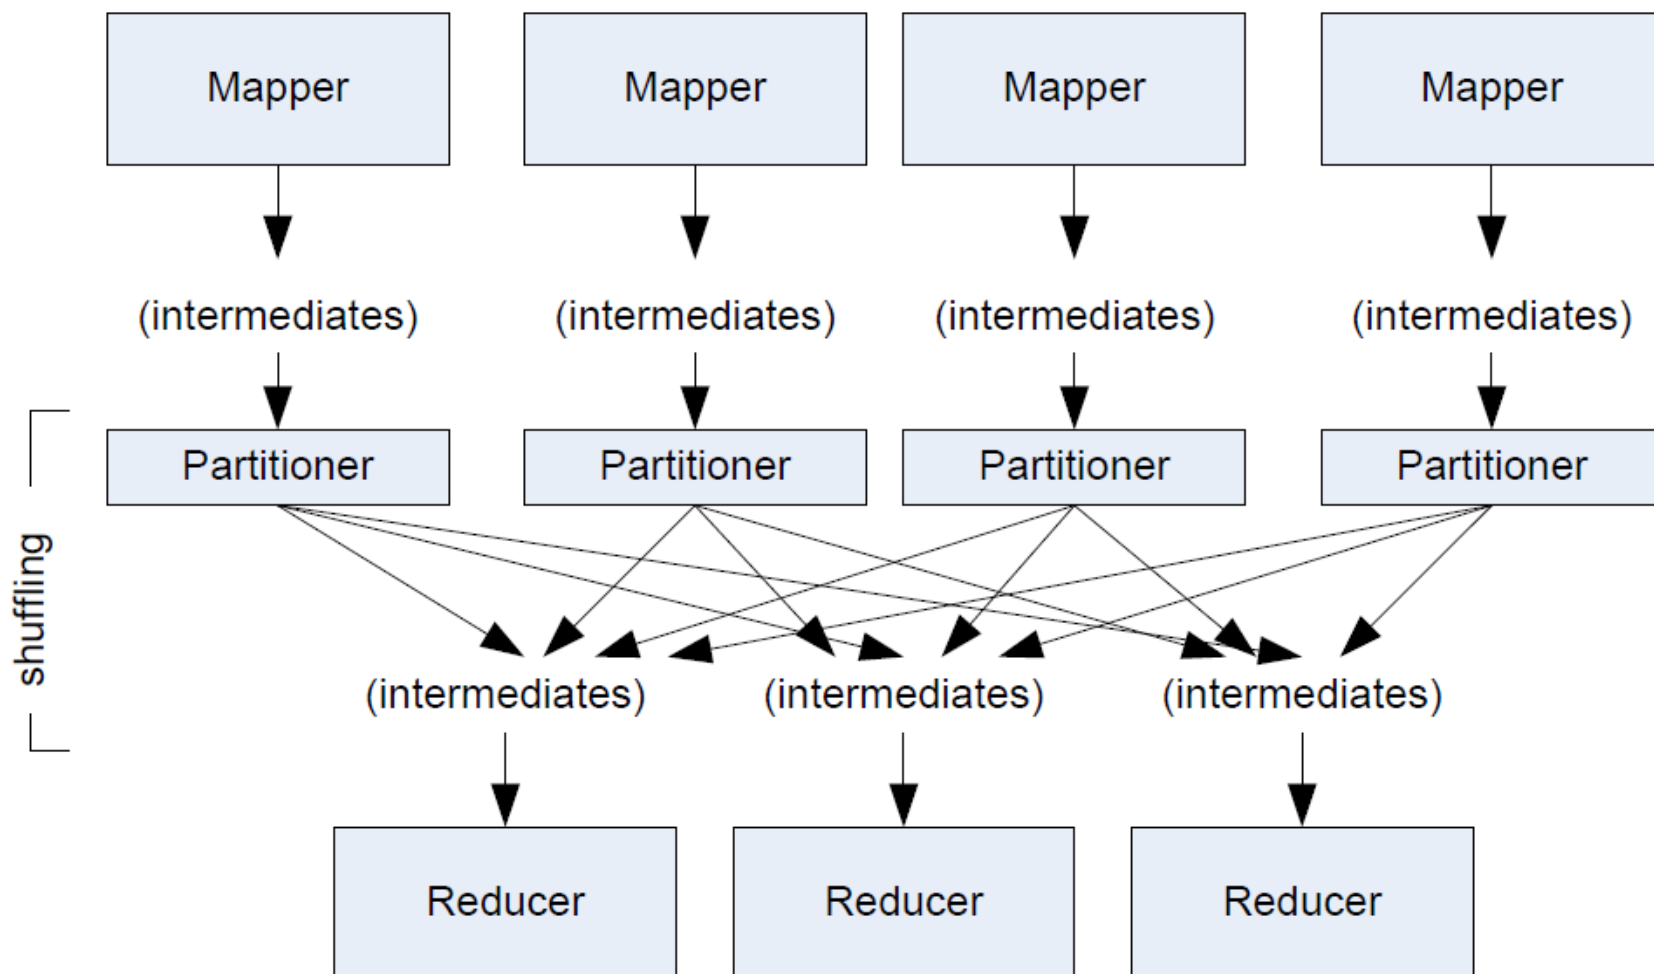
MapReduce Execution Overview

The master distributes M map and R reduce tasks to idle workers

- M == number of shards
- R == the intermediate key space is divided into R parts




Partition and Shuffle




MapReduce Execution Overview

- Each map-task worker reads assigned input shard and outputs intermediate key/value pairs
 - ▣ Output buffered in RAM




MapReduce Execution Overview

- Each worker flushes intermediate values, partitioned into R regions, to disk and notifies the Master process.




MapReduce Execution Overview

- Master process gives disk locations to an available reduce-task worker who reads all associated intermediate data.



MapReduce Execution Overview

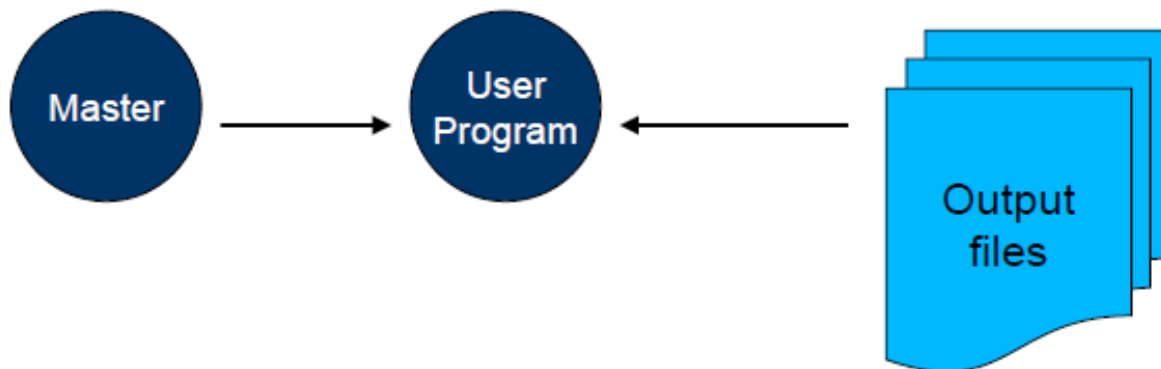
- Each reduce-task worker sorts its intermediate data. Calls the reduce function, passing in unique keys and associated key values. Reduce function output appended to reduce-task's partition output file



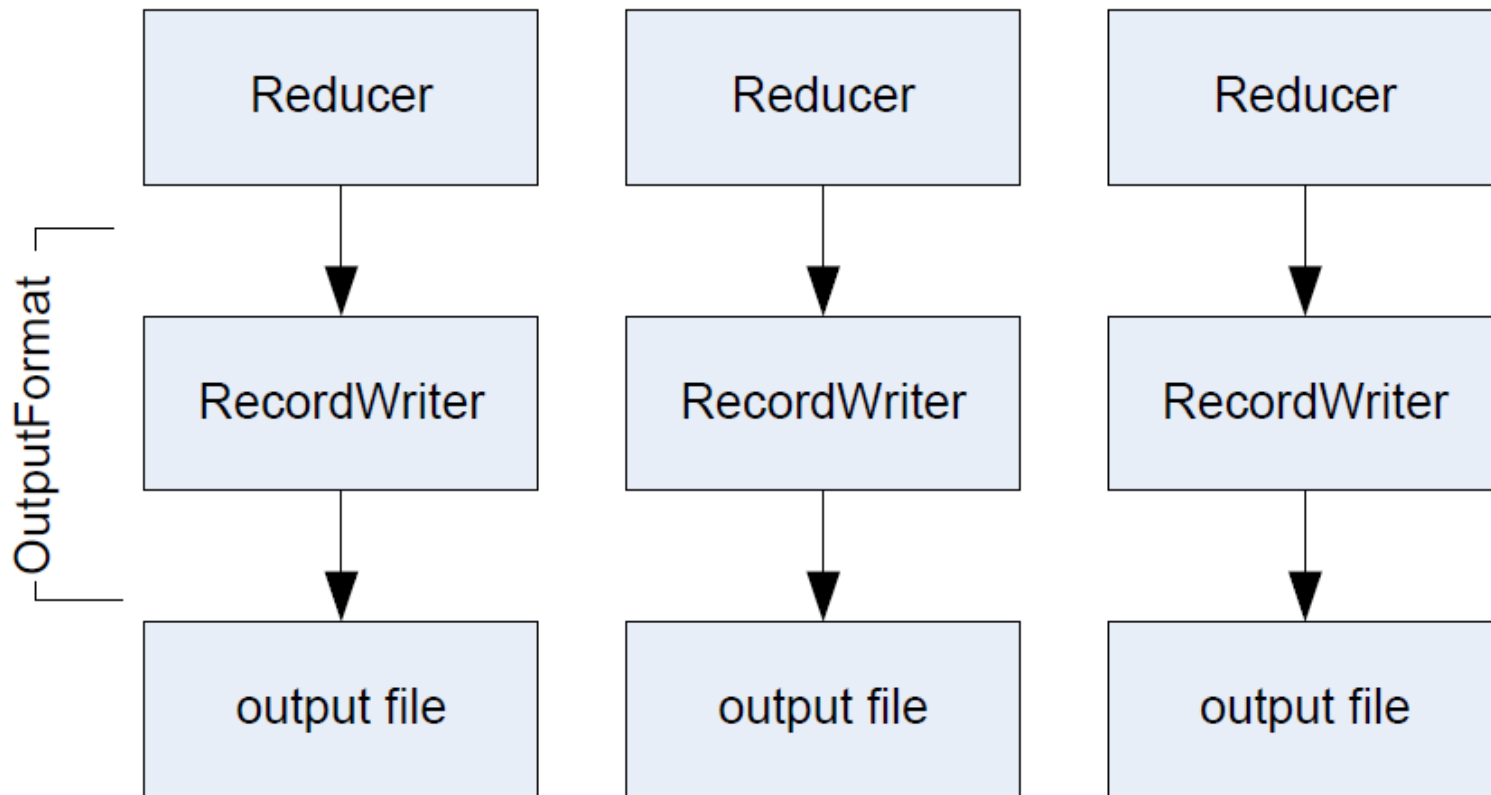
MapReduce Execution Overview

- Master process wakes up user process when all tasks have completed.

Output contained in R output files



Writing The Output



MapReduce Execution Overview

□ Fault Tolerance

- Master process periodically pings workers
 - Map-task failure
 - ✓ Re-execute
 - All output was stored locally
 - Reduce-task failure
 - ✓ Only re-execute partially completed tasks
 - All output stored in the global file system



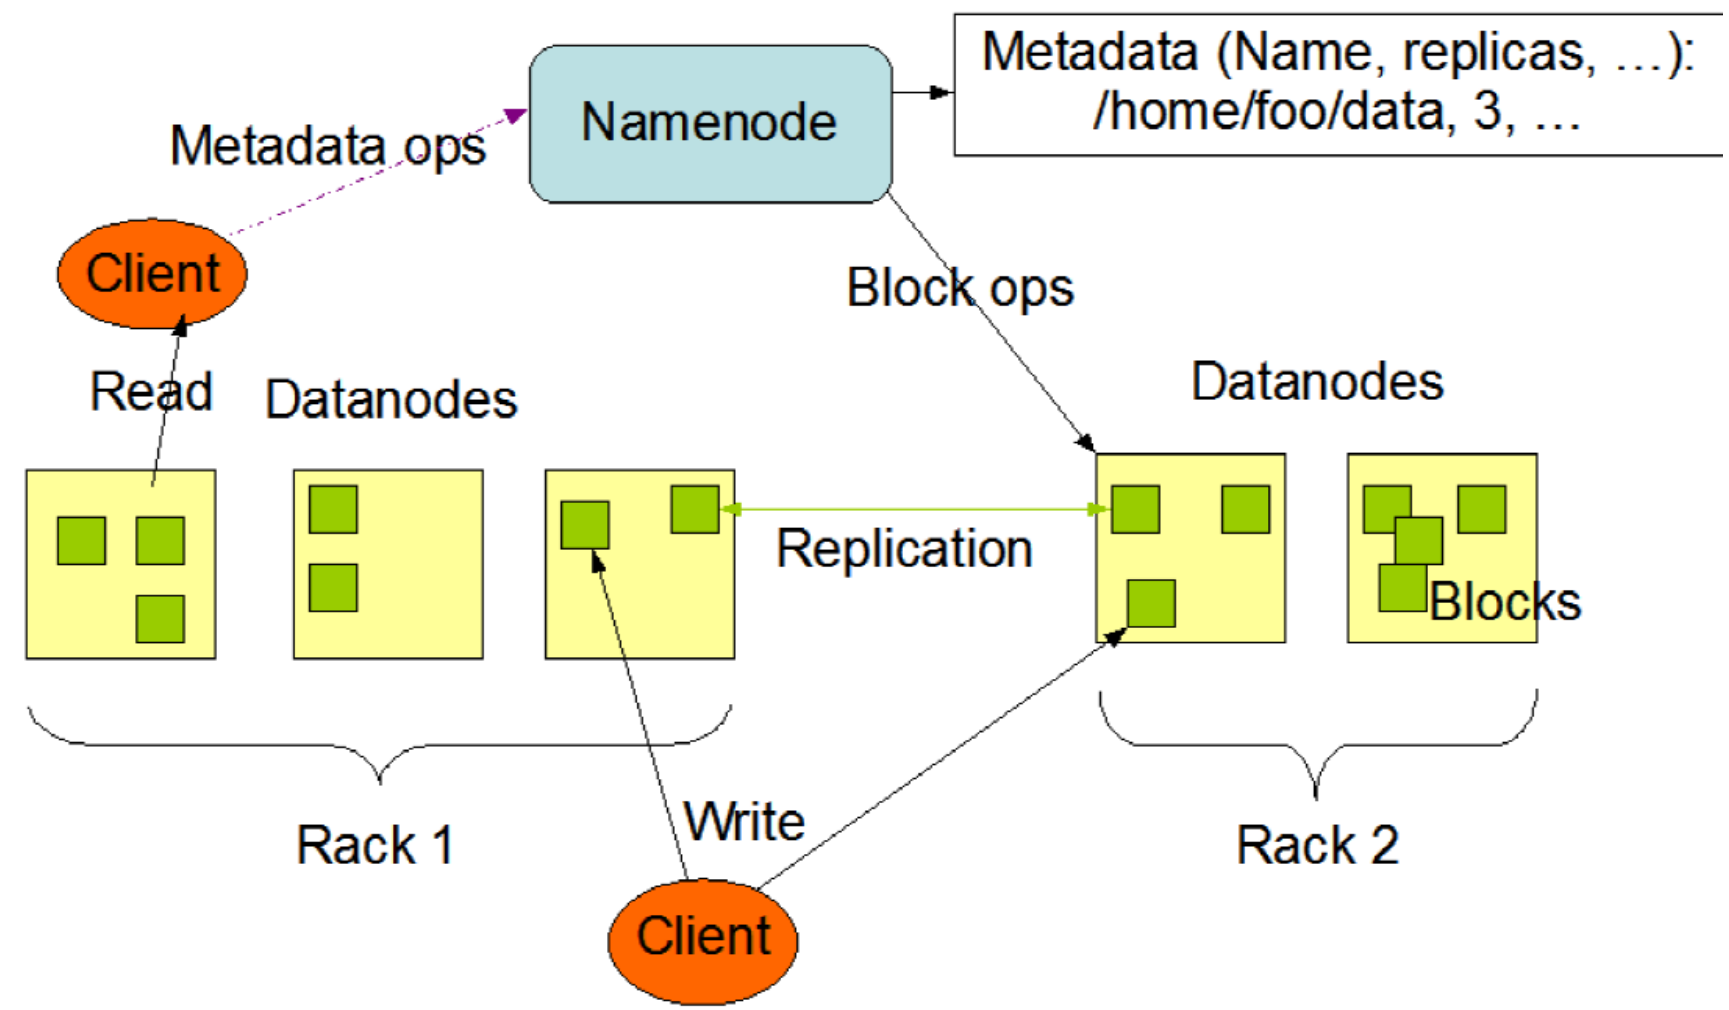
Hadoop

MapReduce Execution Overview

- ❑ Open source MapReduce implementation
 - <http://hadoop.apache.org/core/index.html>

Google calls it	Hadoop equivalent
MapReduce	Hadoop
GFS	HDFS
Bigtable	HBase
Chubby	(nothing yet... but planned)

HDFS Architecture



Hadoop Related Projects

- **Ambari:** A web-based tool for provisioning, managing, and monitoring Apache Hadoop clusters which includes support for Hadoop HDFS, Hadoop MapReduce Hive, HCatalog, HBase, ZooKeeper, Oozie, Pig and Sqoop. Ambari also provides a dashboard for viewing cluster health such as heat maps and ability to view MapReduce, Pig and Hive applications visually along with features to diagnose their performance characteristics in a user-friendly manner
- **Avro:** A data serialization system
- **Cassandra:** A scalable multi-master database with no single points of failure
- **Chukwa:** A data collection system for managing large distributed systems
- **HBase:** A scalable, distributed database that supports structured data storage for large tables (NoSQL)
- **Hive:** A data warehouse infrastructure that provides data summarization and ad hoc querying
- **Mahout:** A Scalable machine learning and data mining library
- **Pig:** A high-level data-flow language and execution framework for parallel computation
- **ZooKeeper:** A high-performance coordination service for distributed applications



References

- Introduction to Parallel Programming and MapReduce, Google Code University
 - ❑ <http://code.google.com/edu/parallel/mapreduce-tutorial.html>
- Distributed Systems
 - ❑ <http://code.google.com/edu/parallel/index.html>
- MapReduce:SimplifiedDataProcessing onLargeClusters
 - ❑ <http://labs.google.com/papers/mapreduce.html>
- Hadoop
 - ❑ <http://hadoop.apache.org/core/>

Thank You !